

Week 10

Assignment Project Exam Help

FIT5202 Big Data Processing

<https://powcoder.com>

Revised by Chee-Ming Ting
(18 May 2021)

Add WeChat powcoder

Data Streaming using Apache Kafka and Spark

Spark Structured Streaming



Week 10 Agenda

- Week 10 Review
 - Apache Kafka
 - Kafka Producer and Kafka Consumer
- Spark Structured Streaming
 - Introduction
 - Typical Use Case with Kafka
 - DEMO :
 - Word Count (Reading from Socket)
 - Click Stream Analysis (Producer and Consumer)
 - **Use case : Log Analysis**
- Other Topics on Structured Streaming
 - Output Modes
 - Output Sink
 - Triggers


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Spark Streaming

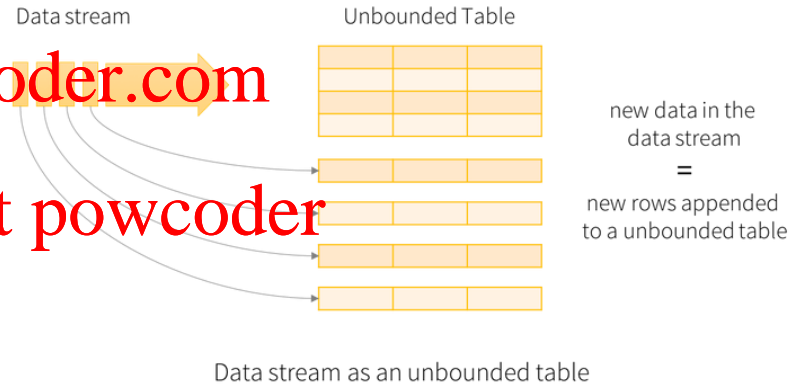
Spark Streaming : used the concept of microbatches, incoming record was a dstream, and RDD based API



Dstream - provide us data divided into chunks as RDDs

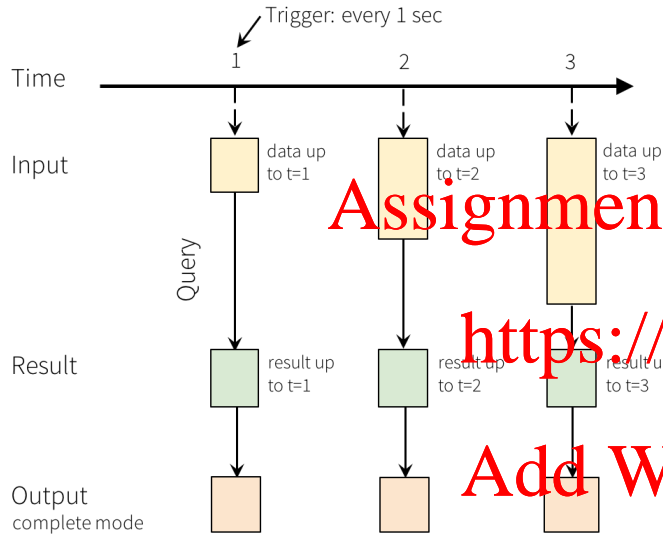
Resource : [Comparison between Spark Streaming and Structured Streaming](#)

Structured Streaming : no concept of batch, data is received in a trigger, appended continuously to the unbounded result table.



Treat a live data stream as a table that is being continuously appended.

Spark Structured Streaming




Programming Model for Structured Streaming


- A query on the input generates the “Result table”
 - At every **trigger** interval (say, every 1 second), new rows are appended to the **input table**, which eventually updates the **result table**
 - Whenever the result table is updated, the changed result rows are written to an external **sink**.
- The output is what is written to external storage, it has 3 modes

1. Complete Mode
2. Append Mode
3. Update Mode


Structured Stream Example (Word Count)



lines streaming dataframe - an unbounded table containing streaming text data

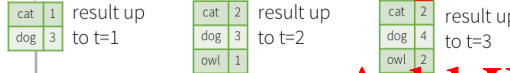


Input
Unbounded
table of all input



word count query

Result
Table of
word counts



Output
Complete Mode



Model of the Quick Example

When **query** is started

- Spark will continuously check for new data
- Spark will run incremental query to combine previous running counts with new data

<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```
# Create DataFrame representing the stream of input lines
lines = spark \
    .readStream \
    .format("socket") \
    .option("host", "localhost") \
    .option("port", 9999) \
    .load()
```

Readstream read
data into DF

```
# Split the lines into words
words = lines.select(
    explode(
        split(lines.value, " ")
    ).alias("word")
)
```

```
# Generate running word count
wordCounts = words.groupBy("word").count()
```

```
# Start running the query that prints the running counts to the console
query = wordCounts \
    .writeStream \
    .outputMode("complete") \
    .format("console") \
    .start()

query.awaitTermination()
```

Clickstream analysis

```
schema = StructType([\n    StructField('Clicks', IntegerType(), True),\n    StructField('Impressions', IntegerType(), True),\n    StructField('ts', TimestampType(), True)\n])
```

Using the schema,
we convert the
data to a Spark
DataFrame

```
df = df.selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)")
```

```
df=df.select(F.from_json(F.col("value").cast("string"), schema).alias('parsed_value'))
```

Cast to String

Parse the string data in json
format according to the schema

```
df_formatted.printSchema()\nroot\n |-- Clicks: integer (nullable = true)\n |-- Impressions: integer (nullable = true)\n |-- ts: timestamp (nullable = true)
```

Assignment Project Exam Help

| key | value | topic | partition | offset | timestamp |
|----------|----------|---------|-----------|--------|------------|
| [binary] | [binary] | "topic" | 0 | 345 | 1486007873 |
| [binary] | [binary] | "topic" | 3 | 389 | 1486006721 |

Kafka format data

- Data in value as byte array

```
root\n |-- key: binary (nullable = true)\n |-- value: binary (nullable = true)\n |-- topic: string (nullable = true)\n |-- partition: integer (nullable = true)\n |-- offset: long (nullable = true)\n |-- timestamp: timestamp (nullable = true)\n |-- timestampType: integer (nullable = true)
```

| Clicks | Impressions | ts |
|--------|-------------|------------|
| 0 | 6 | 1602944650 |
| 0 | 10 | 1602944650 |
| 0 | 5 | 1602944650 |

Query on Streaming Dataframes

```
df = spark \n    .readStream \n    .format("kafka") \n    .option("kafka.bootstrap.servers", "127.0.0.1:9092") \n    .option("subscribe", topic) \n    .load()\nquery = df.trigger()
```

Spark subscribe to
topic & read data

```
#Using the .minute function, we can perform the following aggregation\ngrouped_by_min = df_formatted.groupBy(F.minute("ts").alias("minute_bin"))\n    .agg(F.sum("Impressions").alias("Total Impressions"))
```

Output result table

| minute_bin | Total Impressions |
|------------|-------------------|
| 57 | 172 |
| 54 | 134 |
| 55 | 314 |
| 56 | 290 |

For aggregation query,
use 'complete' mode

```
query = grouped_by_min \n    .writeStream \n    .outputMode("complete") \n    .format("console") \n    .trigger(processingTime='5 seconds') \n    .start()
```

Creating streaming DataFrames

Input Sources

```
topic = "clickstream"
df = spark \
    .readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "127.0.0.1:9092") \
    .option("subscribe", topic) \
    .load()
```

```
lines = spark \
    .readStream \
    .format("socket") \
    .option("host", "localhost") \
    .option("port", 9999) \
    .load()
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

File source - Reads files written in a directory as a stream of data. Files will be processed in the order of file modification time. If latestFirst is set, order will be reversed. Supported file formats are text, CSV, JSON, ORC, Parquet.

Kafka source - Reads data from Kafka. It's compatible with Kafka broker versions 0.10.0 or higher. See the [Kafka Integration Guide](#) for more details.

Socket source (for testing) - Reads UTF8 text data from a socket connection. The listening server socket is at the driver.

Output Modes

- ❑ The output mode specifies **the way the data in result table is written to output sink**

• **Append mode (default)** - Only the new rows added to the Result Table since the last trigger will be outputted to the sink. This is supported for only those queries where rows added to the Result Table is never going to change. Hence, this mode guarantees that each row will be output only once (assuming fault-tolerant sink). For example, queries with only select, where, map, flatMap, filter, join, etc. will support Append mode.



Just write new rows in result table to sink

• **Complete mode** - The whole Result Table will be outputted to the sink after every trigger. This is supported for aggregation queries (e.g., groupBy).



Write all the rows to sink

• **Update mode** - Only the rows in the Result Table that were updated since the last trigger will be outputted to the sink. When there are no aggregations it works exactly the same as "append" mode



Write 'only' all the rows that are updated to sink

```
query = df_formatted \
  .writeStream \
  .outputMode("append") \
  .format("console") \
  .trigger(processingTime='5 seconds') \
  .start()
```

```
query = grouped_by_min \
  .writeStream \
  .outputMode("complete") \
  .format("console") \
  .trigger(processingTime='5 seconds') \
  .start()
```


Output Sink

```
query = df_formatted \
    .writeStream \
    .outputMode("update") \
    .format("console") \
    .start()
```

```
#Change the output sink to "memory" and write output to the memory sink
query = grouped_by_min \
    .writeStream \
    .outputMode("complete") \
    .format("memory") \
    .queryName("impressions_minute_bin") \
    .trigger(processingTime='5 seconds') \
    .start()
```

```
spark.sql("select * from impressions_minute_bin").show()
```

- **File sink** - Stores the output to a directory.

```
writeStream
    .format("parquet") // can be "orc", "json", "csv", etc.
    .option("path", "path/to/destination/dir")
    .start()
```

[What is Parquet file format \[Ref Link\]?](#)

Efficient as well as performant flat columnar storage format compared to row-based csv or tsv files

- **Console sink (for debugging)** - Prints the output to the console/stdout every time there is a trigger. Both, Append and Complete output modes, are supported. This should be used for debugging purposes on low data volumes as the entire output is collected and stored in the driver's memory after every trigger.

```
writeStream
    .format("console")
    .start()
```

- **Memory sink (for debugging)** - The output is stored in memory as an in-memory table. Both, Append and Complete output modes, are supported. This should be used for debugging purposes on low data volumes as the entire output is collected and stored in the driver's memory. Hence, use it with caution.

```
writeStream
    .format("memory")
    .queryName("tableName")
    .start()
```

- **Foreach sink** - Runs arbitrary computation on the records in the output. See later in the section for more details.

```
writeStream
    .foreach(...)
    .start()
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Triggers — How Frequently to Check Sources For New Data

→ defines how often a streaming query should be executed (triggered) and emit the output

```
query = df_formatted |\n    .writeStream |\n    .outputMode("append") |\n    .format("console") |\n    .trigger(processingTime='5 seconds') |\n    .start()
```

1. **Once** : only processes once and terminates the stream

2. **Processing time** : Most widely used and recommended

- ☐ Processes Datastreams as a series of small batch jobs

- ☐ gives better control over how often micro batch jobs should get triggered (e.g. every 5 secs).

3. **Continuous** : Experimental Feature, allows to process records in milliseconds latency

Assignment Project Exam Help

<https://powcoder.com>

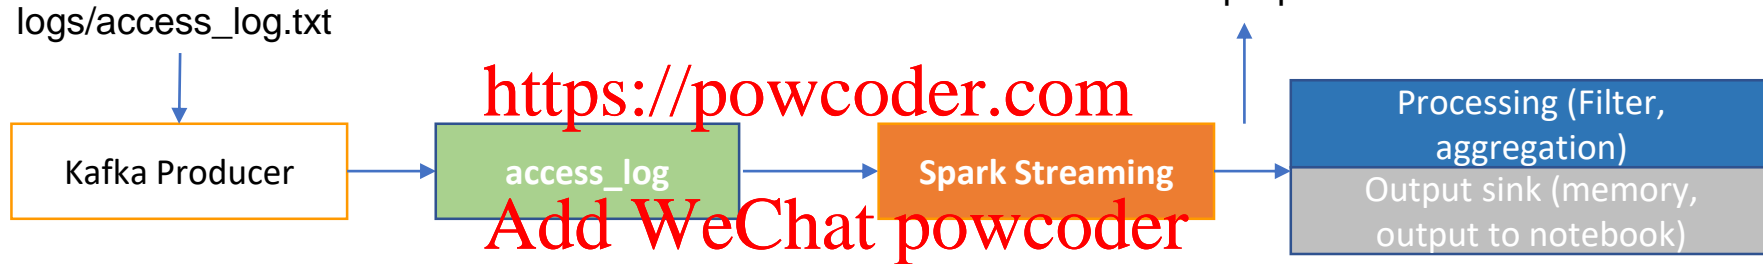
Add WeChat powcoder

Lab Task: Tracking Server Access Log

- ❑ Server is going to continuously send a records of a host who is trying to access some endpoint (url) from the web server
- ❑ **Goal:** Perform real time queries from this data stream and output the results.

Assignment Project Exam Help

Data preparation



<https://powcoder.com>
Add WeChat powcoder

Each line contains some valuable information such as:

1. Host
2. Timestamp
3. HTTP method
4. URL endpoint
5. Status code
6. Protocol
7. Content Size

Task 2: filters those requests that were not successful (status != 200)

Task 3: Count the number of requests by access status code

Task 4: Output the unsuccessful requests (unsucess_df) to memory sink

Handling Array Json

1. Clickstream Spark Streaming - Handling Json Array DEMO

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

References

<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>

<https://docs.microsoft.com/en-us/azure/databricks/getting-started/spark/streaming>

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Thank You!

See you next week.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder