

# Week 11

Assignment Project Exam Help

FIT5202 Big Data Processing

<https://powcoder.com>

Add WeChat powcoder

Data Streaming using Apache Kafka and Spark

Spark Structured Streaming

Aggregations on Windows over event-time

Handling Late events with water marking

# Week 11 Agenda

- Week 10 Review

- Structured Streaming
- Integration with Kafka
- DEMO

- This Week:

- Aggregations on Windows over Event Time
- Handling late data with watermarking
- “Parquet” sink
- Checkpointing

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

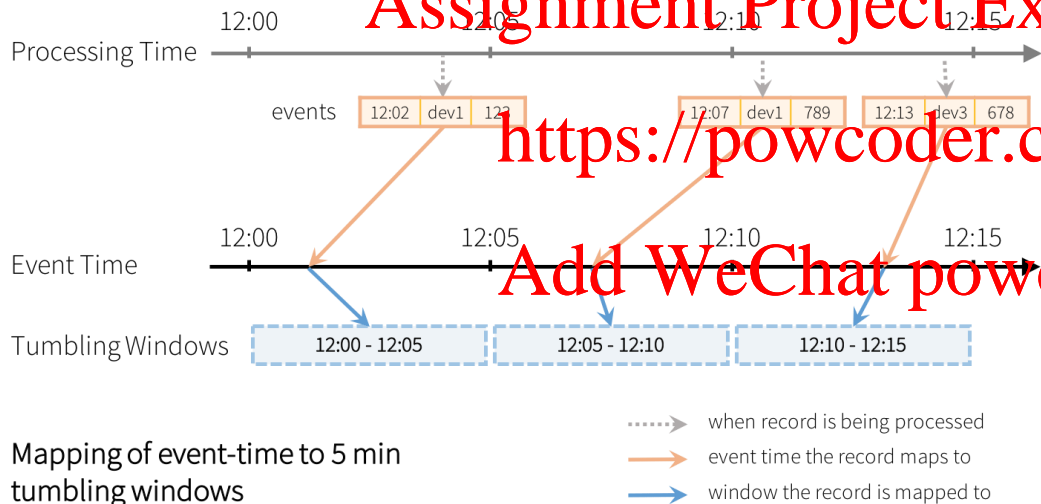
# Aggregations on Windows over Event Time

- ❑ In many cases (e.g., moving averaging), we want aggregations over data bucketed by time windows (e.g., every 5 minutes) rather than over entire streams
- ❑ Bucketing data into windows based on event-time (e.g. the time data generated in the producer) – grouping using **window function**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

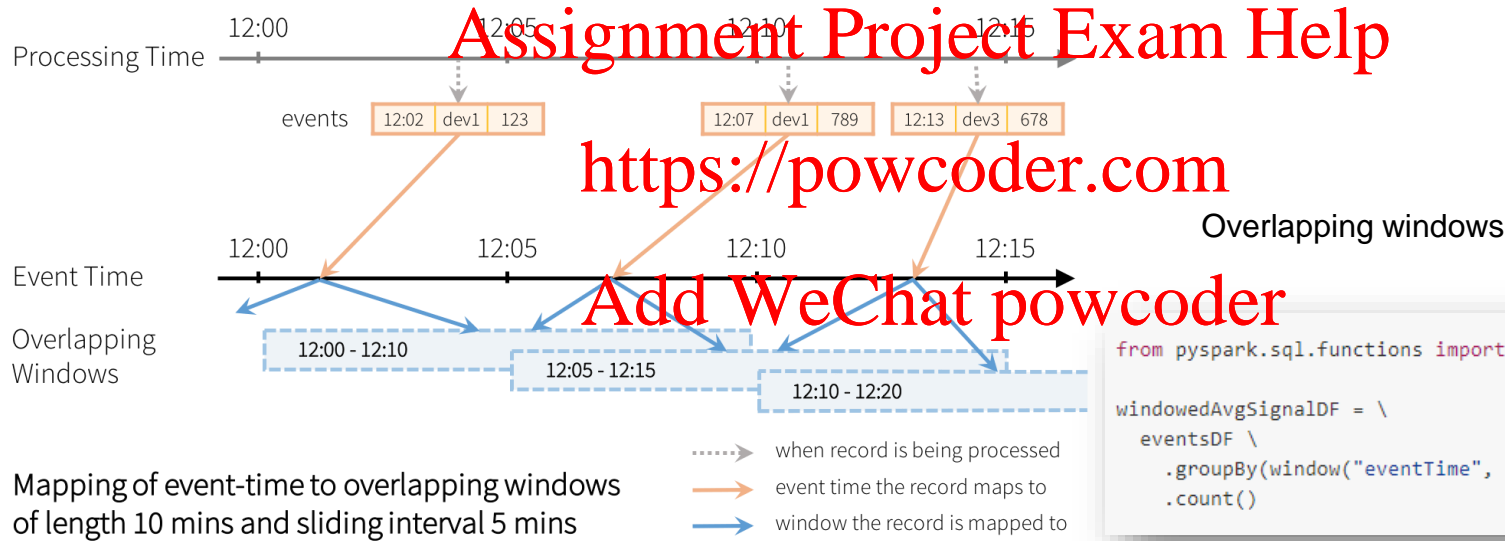


Non-overlapping windows

```
from pyspark.sql.functions import *



windowedAvgSignalDF = \
    eventsDF \
        .groupBy(window("eventTime", "5 minute")) \
        .count()
```

# Aggregations on Windows over Event Time



```
from pyspark.sql.functions import *  
  
windowedAvgSignalDF = \  
    eventsDF \  
        .groupBy(window("eventTime", "10 minutes", "5 minutes")) \  
        .count()
```

# Stateless vs Stateful Stream Processing



Assignment Project Exam Help

<https://powcoder.com>

Stateful

Add WeChat powcoder

## Stateless

- ❑ Each record is processed independently of other records
- ❑ e.g. operations like map, filter, join with static data

- ❑ Processing of records depends upon the result of previously processed records.

- ❑ Need to maintain “**intermediate information**” for processing – called “**state**”
- ❑ E.g., operations like aggregating count of records (e.g., intermediate count)

## State of Progress

- ❑ keeping track of data that has been processed in streaming so far.
- ❑ Called **checkpointing**/saving of offsets of incoming data.

## State of Data

- ❑ intermediate information derived from data (processed so far).

# Watermarking

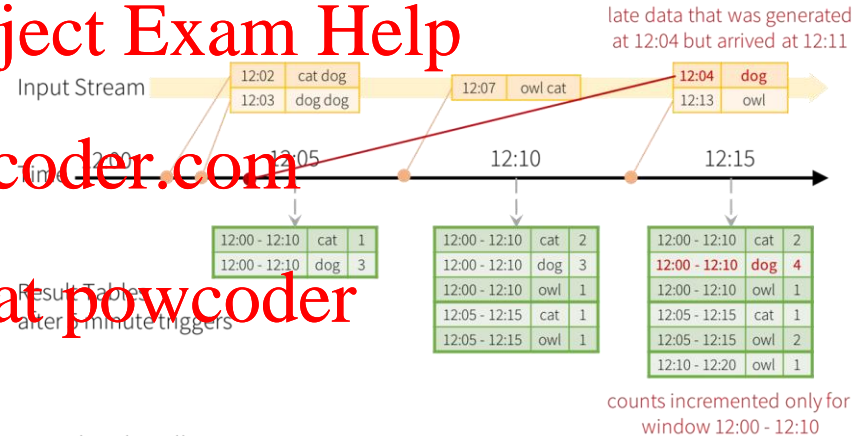
Structured Streaming can maintain the intermediate state for partial aggregates (e.g. intermediate count) for a period of time such that late data can update aggregates of old windows correctly

- To handle events that arrive late to the application
- E.g. a word is generated at 12:04 (event time) but received at 12:11 by the application.
- The application should use the time 12:04 instead of 12:11 to update the older counts for a window 12:00 - 12:10.
- Watermarking lets engine automatically track the current event time in data and clean up/update old state accordingly

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Late data handling in  
Windowed Grouped Aggregation

# Watermarking

Two parameters to define the watermark of a query

- (1) **event time column**
- (2) **Threshold** specify for how late data should be processed (in event time)



Assignment Project Exam Help

```
withWatermark(eventTime: String, delayThreshold: String)
```

late data within the threshold will be aggregated, but data later than the threshold will start getting dropped

**max event time** - is the latest event time seen by the engine

<https://powcoder.com>


Add WeChat powcoder

```
# Group the data by window and word and compute the count of each group
windowedCounts = words \
    .withWatermark("timestamp", "10 minutes") \
    .groupBy(
        window(words.timestamp, "10 minutes", "5 minutes"),
        words.word) \
    .count()
```

# Watermarking

Window-based aggregation  
based on event time –  
Window size = 10 mins  
Slide = 5 mins

```
# Group the data by window and word and compute the count of each group
windowedCounts = words \
    .withWatermark("timestamp", "10 minutes") \
    .groupBy(
        window(words.timestamp, "10 minutes", "5 minutes"),
        words.word) \
    .count()
```



Processing Time with 5 min triggers

Result Tables after each trigger

Watermarking in Windowed  
Grouped Aggregation with Update Mode

12:00 - 12:10	owl	1
12:00 - 12:10	dog	1
12:05 - 12:15	owl	1
12:05 - 12:15	dog	1

12:05 - 12:15	owl	1
12:05 - 12:15	dog	2
12:05 - 12:15	cat	1
12:10 - 12:20	dog	1

12:00 - 12:10	dog	2
12:05 - 12:15	owl	2
12:05 - 12:15	dog	3
12:05 - 12:15	cat	2
12:10 - 12:20	dog	1
12:10 - 12:20	cat	1
12:10 - 12:20	owl	1

12:00 - 12:10	dog	2
12:05 - 12:15	owl	2
12:05 - 12:15	dog	3
12:05 - 12:15	cat	2
12:10 - 12:20	dog	1
12:10 - 12:20	cat	1
12:10 - 12:20	owl	2

table updated with too late data (12:04, donkey)

table updated with late data (12:17, owl)

purple rows are updated rows that are written to the sink as output

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Engine will keep updating counts of a window in the Result Table until the window is older than the watermark, which lags behind the current event time by 10 minutes.



# Recovering from Failures with Checkpointing

- In case of failure, can recover the previous progress and state of previous query and continue where it left off.
- Can enable checkpointing using the option **checkpointLocation** on the query.
- To save all the progress information (i.e. range of offsets processed in each trigger) and the running aggregates ('states') to the checkpoint location

```
golf \
  .writeStream \
  .outputMode("complete") \
  .option("checkpointLocation", "path/to/HDFS/dir") \
  .asTable() \
  .start()
```

# Clickstream Watermarking

```
schema = StructType([\n    StructField('Clicks', IntegerType(), True),\n    StructField('Impressions', IntegerType(), True),\n    StructField('ts', TimestampType(), True)\n])
```

Using the schema,  
we convert the  
data to a Spark  
DataFrame

```
df=df.select(F.from_json(F.col("value").cast("string"), schema).alias('parsed_value'))
```

```
df = df.selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)")
```

Kafka producer  
Late events

```
-----\n2021-05-20 23:46:15 impressions: 7\n2021-05-20 23:46:15 impressions: 9\n2021-05-20 23:46:15 impressions: 4\n2021-05-20 23:46:15 impressions: 2\n2021-05-20 23:46:15 impressions: 1\n2021-05-20 23:46:15 impressions: 7\n2021-05-20 23:46:15 impressions: 3\n2021-05-20 23:46:15 impressions: 7\n2021-05-20 23:46:15 impressions: 4\n2021-05-20 23:46:10 impressions: 8\n2021-05-20 23:46:10 impressions: 2\n2021-05-20 23:46:10 impressions: 5\n-----\n2021-05-20 23:46:20 impressions: 5\n2021-05-20 23:46:20 impressions: 1\n2021-05-20 23:46:20 impressions: 4\n-----
```

Cast to String

Parse the string data in json  
format according to the schema

Assignment Project Exam Help

key	value	topic	partition	offset	timestamp
[binary]	[binary]	"topic"	0	345	1486007873
[binary]	[binary]	"topic"	3	389	1486007873

Spark subscribe to topic &  
read data

```
root\n-- key: binary (nullable = true)\n-- value: binary (nullable = true)\n-- topic: string (nullable = true)\n-- partition: integer (nullable = true)\n-- offset: long (nullable = true)\n-- timestamp: timestamp (nullable = true)\n-- timestampType: integer (nullable = true)
```

```
df_formatted.printSchema()\nroot\n-- Clicks: integer (nullable = true)\n-- Impressions: integer (nullable = true)\n-- ts: timestamp (nullable = true)
```

Clicks	Impressions	ts
0	6	1602944650
0	10	1602944650
0	5	1602944650

Aggregations on Windows over  
Event-time

```
windowedCounts = df_formatted \n    .withWatermark("ts", "10 seconds") \n    .groupBy(window(df_formatted.ts, "10 seconds")) \n    .agg(F.sum("Impressions").alias("total")) \n    .select("window", "total")
```

Output result table

```
query = windowedCounts \n    .writeStream \n    .outputMode("complete") \n    .format("console") \n    .trigger(processingTime='5 seconds') \n    .option("truncate", "false") \n    .start()
```

window	total
[2021-05-20 23:52:00, 2021-05-20 23:52:10]	124
[2021-05-20 23:51:40, 2021-05-20 23:51:50]	138
[2021-05-20 23:52:20, 2021-05-20 23:52:30]	40
[2021-05-20 23:51:30, 2021-05-20 23:51:40]	85
[2021-05-20 23:51:50, 2021-05-20 23:52:00]	134
[2021-05-20 23:52:10, 2021-05-20 23:52:20]	133


For aggregation query,  
use 'complete' mode

# Granularity Reduction DEMO

(Refer to Lecture)

The Aggregation on Windows over Event-time is another way of understanding the concept of granularity reduction.

Real-time uniform stream data visualization



Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

# Lab Task: Access Log – Window-based Aggregation

## Aggregations on window over event-time

tsExp=r'(\d{10})\s'  
- Extract event timestamp

logs/access\_log.txt

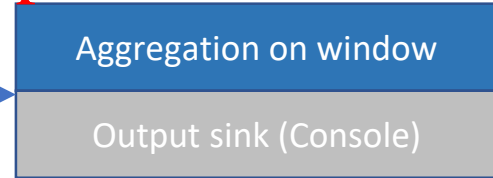


Assignment Project Exam Help

<https://powcoder.com>



Data preparation



Append with event time, ts

Each line contains some valuable information such as:

1. Host
2. Timestamp
3. HTTP method
4. URL endpoint
5. Status code
6. Protocol
7. Content Size

Add WeChat powcoder

Task 1:

- ☐ Using the Window function, find the number of logs for each status in a window of 30 seconds. Set the window slide to 10 seconds
- ☐ Write the output to console sink.

# References

- <https://databricks.com/blog/2017/05/08/event-time-aggregation-watermarking-apache-sparks-structured-streaming.html>
- <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#quick-example>
- <https://docs.databricks.com/spark/latest/structured-streaming/production.html>
- <http://blog.madhukaraphatak.com/introduction-to-spark-structured-streaming-part-7/>

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder