

# FIT5202 (Volume IV – Sort and Group By) Assignment Project Exam Help

Week 4a – Parallel Sort

<https://powcoder.com>

**algorithm** distributed systems **database**  
systems **computation** knowledge management  
**design** e-business **model** data mining **int**  
distributed systems **database** software  
**computation** knowledge management **an**

# Chapter 5

## Parallel Join

TANIA  
LEUNG  
RAHAYU  
GOEL

Wiley Series on Parallel and Distributed Computing • Albert Zomaya, Series Editor

High Performance Parallel  
Database Processing and  
Grid Databases

High Performance Parallel Database  
Processing and Grid Databases

DAVID TANIA, CLEMENT H.C. LEUNG,  
WENNY RAHAYU, and SUSHANT GOEL



WILEY



- 5.1 Join Operations
- 5.2 Serial Join Algorithms
- 5.3 Parallel Join Algorithms
- 5.4 Cost Models
- 5.5 Parallel Join Optimization
- 5.6 Summary
- 5.7 Bibliographical Notes
- 5.8 Exercises

# Revision

- **Exercise 1 (FLUX Quiz)**

- Parallel Join algorithms for Inner Join consists of two major phases: Data Partitioning, and Local Join

Assignment Project Exam Help

- A. TRUE
- B. FALSE

<https://powcoder.com>

Add WeChat powcoder

# Revision

- **Exercise 2 (FLUX Quiz)**

- Parallel Join algorithms for Outer Join queries are:

- A. ROJA and DOJA
- B. DER
- C. OJSO
- E. only A and B
- F. A, B and C are correct.

Assignment Project Exam Help

<https://powcoder.com>

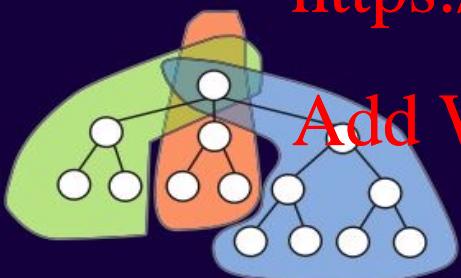
Add WeChat powcoder

TANIA  
LEUNG  
RAHAYU  
GOEL

Wiley Series on Parallel and Distributed Computing • Albert Zomaya, Series Editor

High Performance Parallel  
Database Processing and  
Grid Databases

High Performance Parallel Database  
Processing and Grid Databases



DAVID TANIA, CLEMENT H.C. LEUNG,  
WENNY RAHAYU, and SUSHANT GOEL



WILEY

# Chapter 4

## Parallel Sort and GroupBy

<https://powcoder.com>

Add WeChat powcoder

- 4.1 Sorting, Duplicate Removal and Aggregate
- 4.2 Serial External Sorting Method
- 4.3 Algorithms for Parallel External Sort
- 4.4 Parallel Algorithms for GroupBy Queries
- 4.5 Cost Models for Parallel Sort
- 4.6 Cost Models for Parallel GroupBy
- 4.7 Summary
- 4.8 Bibliographical Notes
- 4.9 Exercises

# 4.1. Sorting, and Serial Sorting

- Serial Sorting – **INTERNAL**
  - The data to be sorted fits entirely into the main memory  
**Assignment Project Exam Help**
- Serial Sorting - **E**[INTERNAL](https://powcoder.com)
  - The data to be sorted DOES NOT fit entirely into the main memory  
**Add WeChat powcoder**

## 4.1. Internal Serial Sorting (cont'd)

- **Bubble Sort**

- Based on swapping
- It compares the first two elements, and if the first is greater than the second, it swaps them.
- It continues doing this for each pair of adjacent elements to the end of the data set.
- It then starts again with the first two elements, repeating until no swaps have occurred on the last pass.
- Example: 6 5 3 1 8 7 2 4

## 4.1. Internal Serial Sorting (cont'd)

6 5 3 1 8 7 2 4

# Bubble Sort

6 5 3 1 8 7 2 4

**5 3 1 6 7 2 4 8**      **3 1 5 6 2 4 7 8**      **1 3 5 2 4 6 7 8**      **1 3 2 4 5 6 7 8**

5 6 3 1 8 7 2 4

3 5 16 7 24 8 Assignment Project Exam Help 13 5 6 24 6 7 8 1 3 2 4 5 6 7 8

5 3 6 1 8 7 2 4

3 1 5 6 7 2 4 8      1 3 5 6 2 4 7 8      1 3 5 2 4 6 7 8      1 2 3 4 5 6 7 8

53168724

3 1 5 6 7 2 4 8 1 3 7 6 9 4 7 8 1 3 2 5 4 6 7 8 1 2 3 4 5 6 7 8

5 3 1 6 8 7 2 4

3 1 5 6 7 2 4 8      1 3 5 2 6 4 7 8      1 3 2 4 5 6 7 8

5 3 1 6 7 8 2 4

3 1 5 6 2 7 Add WeChat nowcoder

5 3 1 6 7 2 8 4

3 1 5 6 2 4 7 8

5 3 1 6 7 2 4 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

1 2 3 4 5 6 7 8

12345678

**Finished**



## 4.1. Internal Serial Sorting (cont'd)

- **Insertion Sort**

6 5 3 1 8 7 2 4

- Based on inserting a new value
- It works by taking elements from the list one by one and inserting them in their correct position into a new sorted list. In arrays, the new list and the remaining elements can share the array's space, but insertion is expensive, requiring shifting all following elements over by one.
- Example: 6 5 3 1 8 <https://powcoder.com>

- 6 5 3 1 8 7 2 4
- 6 5 3 1 8 7 2 4
- 5 6 3 1 8 7 2 4
- 3 5 6 1 8 7 2 4
- 1 3 5 6 8 7 2 4
- 1 3 5 6 8 7 2 4
- 1 3 5 6 7 8 2 4
- 1 2 3 5 6 7 8 4

Add WeChat [powcoder](#)

- Take out 6, and insert it in the previous list
- Take out 5, and insert it in the previous list
- Take out 3, and insert it in the previous list
- Take out 1, and insert it in the previous list
- Take out 8, and insert it in the previous list
- Take out 7, and insert it in the previous list
- Take out 2, and insert it in the previous list
- Take out 4, and insert it in the previous list

6 5 3 1 8 7 2 4

5 6 3 1 8 7 2 4

3 5 6 1 8 7 2 4

1 3 5 6 8 7 2 4

1 3 5 6 8 7 2 4

1 3 5 6 7 8 2 4

1 2 3 5 6 7 8 4

1 2 3 4 5 6 7 8

Finished

## 4.1. Internal Serial Sorting (cont'd)

- **Quick Sort**

- Quick Sort is a divide and conquer algorithm which relies on a partition operation: to partition an array an element called a **pivot** is selected.
- All elements smaller than the pivot are moved before it and all greater elements are moved after it. <https://powcoder.com>
- The lesser and greater sublists are then recursively sorted.
- The most complex issue in Quick Sort is choosing a good pivot element; consistently poor choices of pivots can result in drastically slower performance
- Example: 6 5 3 1 8 7 2 4

6 5 3 1 8 7 2 4

## 4.1. Internal Serial Sorting (cont'd)

- **Quick Sort**

- Quick Sort is a divide and conquer algorithm which relies on a partition operation: to partition an array an element called a **pivot** is selected.
- All elements smaller than the pivot are moved before it and all greater elements are moved after it. <https://powcoder.com>
- The lesser and greater sublists are then recursively sorted.
- The most complex issue in Quick Sort is choosing a good pivot element; consistently poor choices of pivots can result in drastically slower performance
- Example: 6 5 3 1 8 7 2 4

6 5 3 1 8 7 2 4

**Homework: Work out step-by-step  
this Quick Sort example**

## 4.2. Serial External Sorting

- Sorting is expressed by the ORDER BY clause in SQL
- Duplicate remove is identified by the keyword DISTINCT in SQL

Assignment Project Exam Help

**Query 4.1:**

```
Select *  
From STUDENT  
Order By Sdegree;
```

**Query 4.3:**

```
Select Distinct Sdegree  
From STUDENT;
```

<https://powcoder.com>

Add WeChat powcoder

## 4.2. Serial External Sorting (cont'd)

- External sorting assumes that the data does not fit into main memory
- Most common external sorting is sort-merge
- Break the file up into unsorted subfiles, sort the subfiles, and then merge the subfiles into larger and larger sorted subfiles until the entire file is sorted

Assignment Project Exam Help

---

Algorithm: Serial External Sorting

```
// Sort phase - Pass 0
1. Read  $B$  pages at a time into memory
2. Sort them, and Write out a sub-file
3. Repeat steps 1-2 until all pages have been processed

// Merge phase - Pass  $i = 1, 2, \dots$ 
4. While the number of sub-files at end of previous pass
   is  $> 1$ 
5. While there are sub-files to be merged from
   previous pass
6. Choose  $B-1$  sorted sub-files from the previous pass
7. Read each sub-file into an input buffer page
   at a time
8. Merge these sub-files into one bigger sub-file
9. Write to the output buffer one page at a time
```

---

Figure 4.1 External sorting algorithm based on sort-merge

## 4.2. Serial External Sorting (cont'd)

- Example

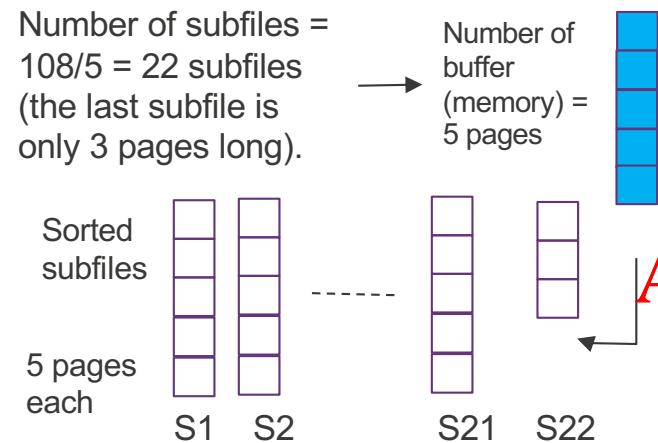
- File size to be sorted = 108 pages, number of buffer (or memory size) = 5 pages
- Number of subfiles =  $108/5 = 22$  subfiles (the last subfile is only 3 pages long).
- **Pass 0** (sorting phase): For each subfile, **read from disk, sort in main-memory**, and **write to disk** (Note: sorting the data in main-memory can use any fast in-memory sorting method, like Quick Sort)

<https://powcoder.com>

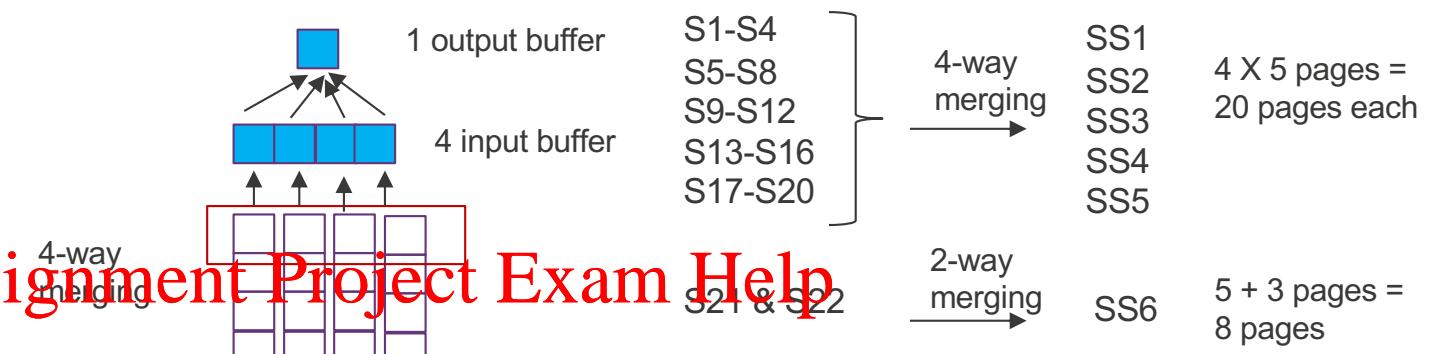
- Merging phase: **We use  $B-1$  buffers (4 buffers) for input and 1 buffer for output**
- **Add WeChat powcoder** **Pass 1:** Read 4 sorted subfiles and perform 4-way merging (apply a need  $k$ -way algorithm). Repeat the 4-way merging until all subfiles are processed. Result = 6 subfiles with 20 pages each (except the last one which has 8 pages)
- **Pass 2:** Repeat 4-way merging of the 6 subfiles like pass 1 above. Result = 2 subfiles
- **Pass 3:** Merge the last 2 subfiles
- Summary: 108 pages and 5 buffer pages require 4 passes



### Pass 0 (sorting phase):



### Pass 1 (Merging phase):



Assignment Project Exam Help

<https://powcoder.com>

### Pass 2 (Merging phase):



### Add WeChat powcoder

### Pass 3 (Merging phase):



## 4.2. Serial External Sorting (cont'd)

### • Exercise 3 (FLUX Quiz)

- There are 150 data pages to be sorted. The machine that we have has a limited memory, and can only take 8 pages at a time. How many passes will it take to sort the 150 data pages?

Assignment Project Exam Help

- A. 2
- B. 3
- C. 4
- D. 5

<https://powcoder.com>

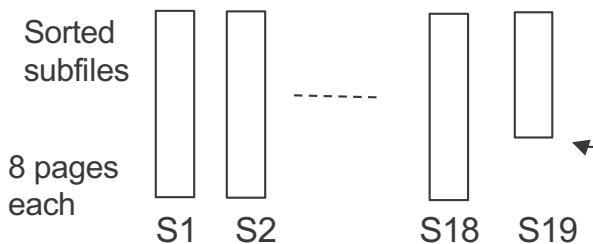
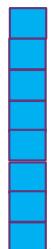
Add WeChat powcoder

File size to be sorted = 150 pages, number of buffer (or memory size) = 8 pages

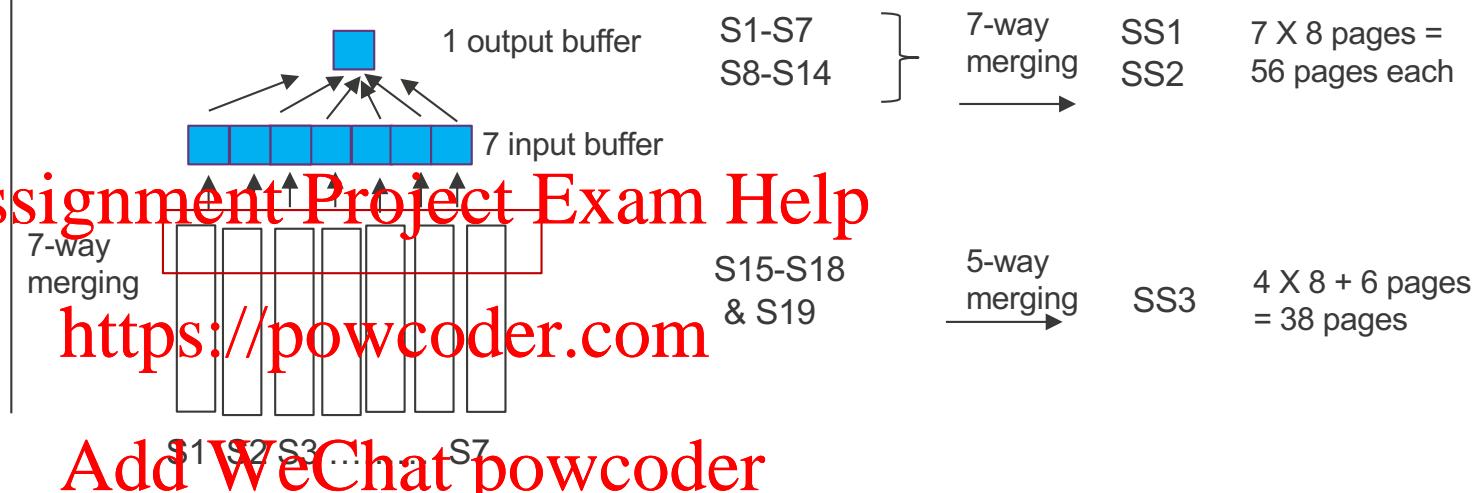
### Pass 0 (sorting phase):

Number of subfiles =  
 $150/8 = 19$  subfiles  
(the last subfile is  
only 6 pages long).

→ Number of  
buffer  
(memory) =  
8 pages



### Pass 1 (Merging phase):



### Pass 2 (Merging phase):



## 4.2. Serial External Sorting (cont'd)

- Example

- Buffer size plays an important role in external sort

Assignment Project Exam Help

**Table 4.1** Number of passes in serial external sorting as number of buffer increases

R	B = 3	B = 5	B = 9	B = 17	B = 129	B = 257
<b>100</b>	2	4	3	2	1	1
<b>1,000</b>	10	5	4	3	2	2
<b>10,000</b>	13	7	5	4	2	2
<b>100,000</b>	17	9	6	5	3	3
<b>1 million</b>	20	10	7	5	3	3
<b>10 million</b>	23	12	8	6	4	3
<b>100 million</b>	26	14	9	7	4	4
<b>1 billion</b>	30	15	10	8	5	4

## 4.3. Parallel External Sort

5 different Algorithms

- Parallel Merge-All Sort
  - Parallel Binary-Merge Sort
  - Parallel Redistribution Binary-Merge Sort
  - Parallel Redistribution Merge-All Sort
  - Parallel Partitioned Sort
- Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder
- 

## 4.3. Parallel External Sort (cont'd)

- **Parallel Merge-All Sort**

- A traditional approach
- Two phases: local sort and final merge
- Load balanced in local sort
- Problems with merging:
  - Heavy load on one processor
  - Network contention

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

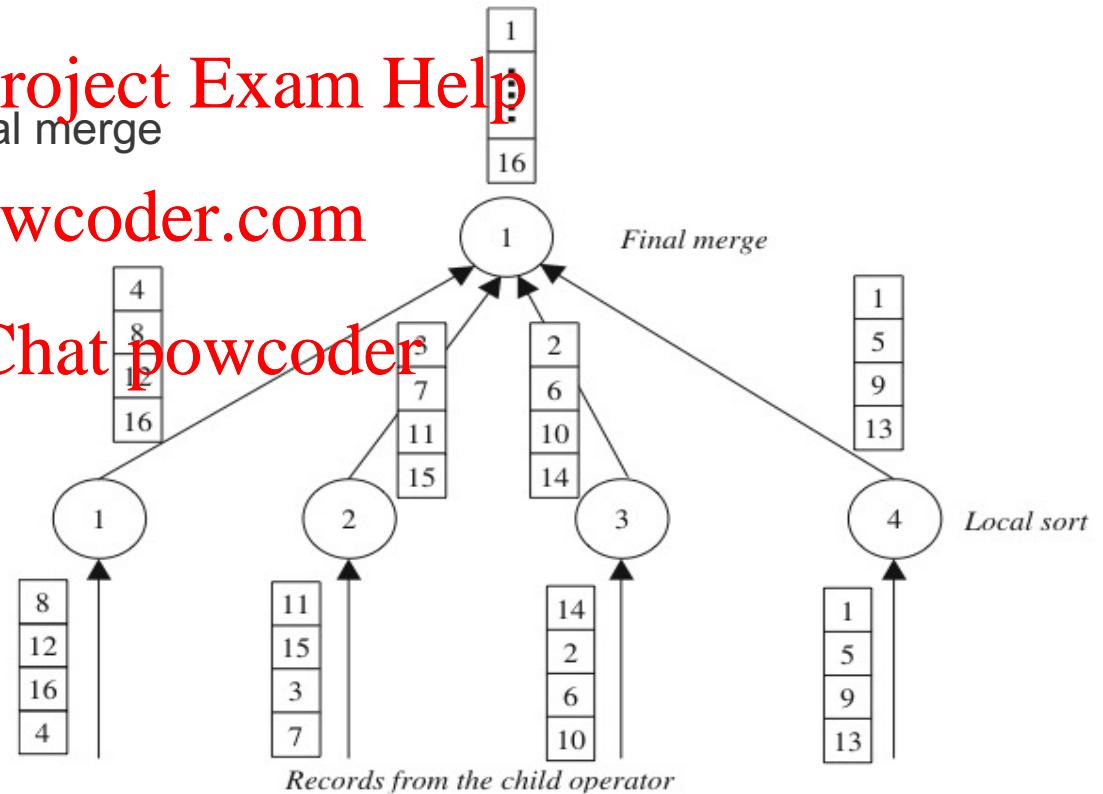


Figure 4.3 Parallel merge-all sort

## 4.3. Parallel External Sort (cont'd)

### • Parallel Binary-Merge Sort

- Local sort similar to traditional method
- Merging in pairs only
- Merging work is now spread to pipeline of processors,  
but merging is still heavy

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

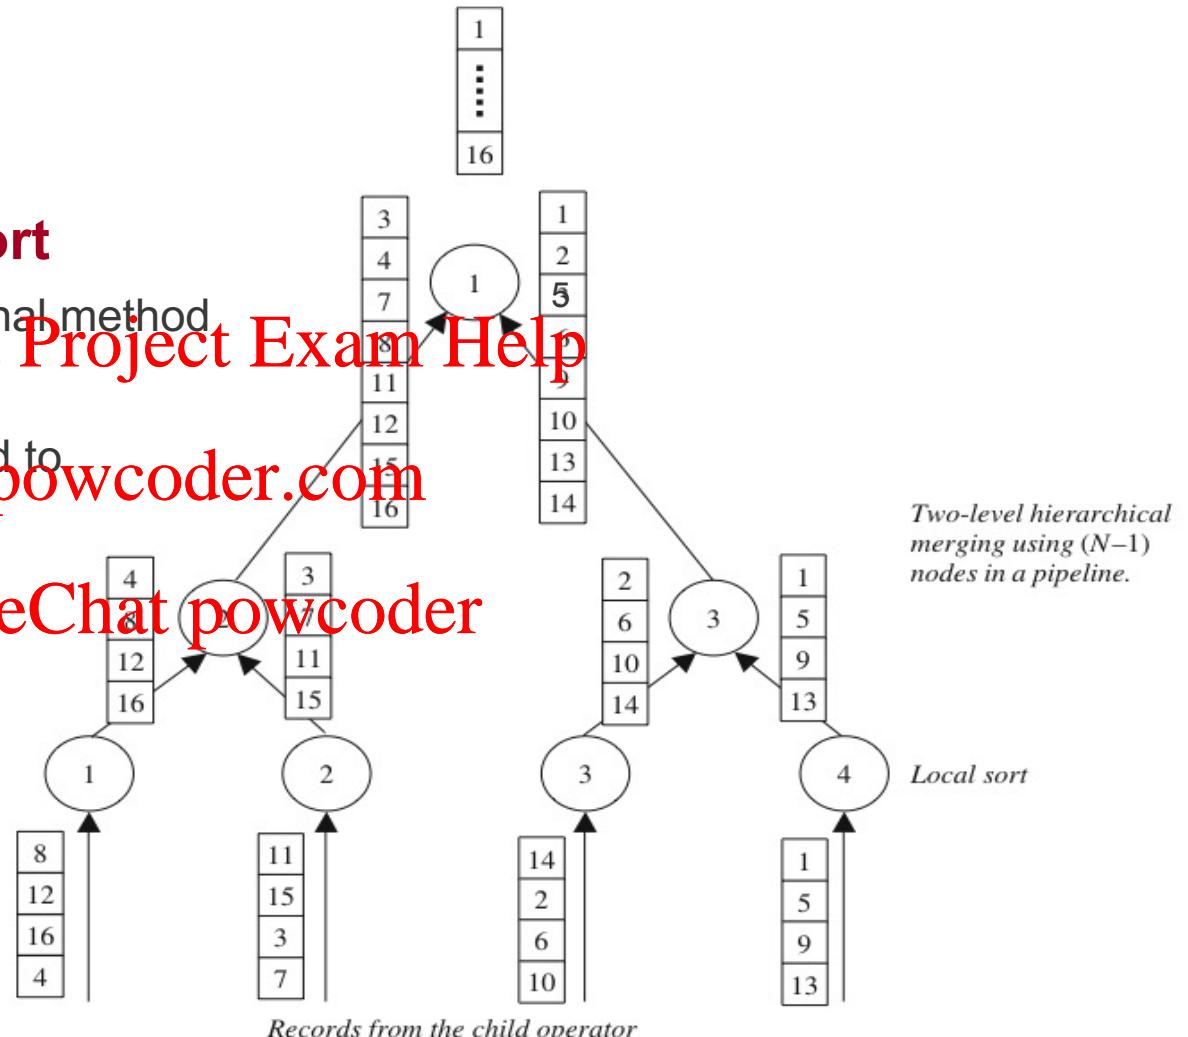


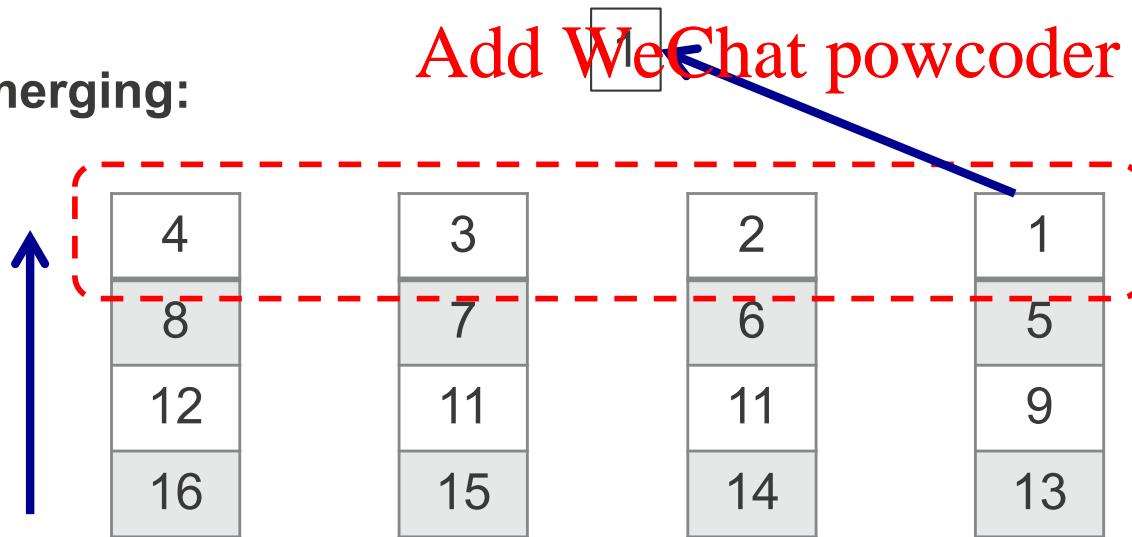
Figure 4.4 Parallel binary-merge sort

## 4.3. Parallel External Sort (cont'd)

- **Parallel Binary-Merge Sort**

- Binary merging vs.  $k$ -way merging
- In  $k$ -way merging, the **searching for the smallest value among  $k$  partitions** is done at the same time
- In binary merging, it is pairwise, but can be time consuming if the list is long
- System requirements:  **$k$ -way merging requires  $k$  files open simultaneously**, but the pipeline process in binary merging requires extra overheads

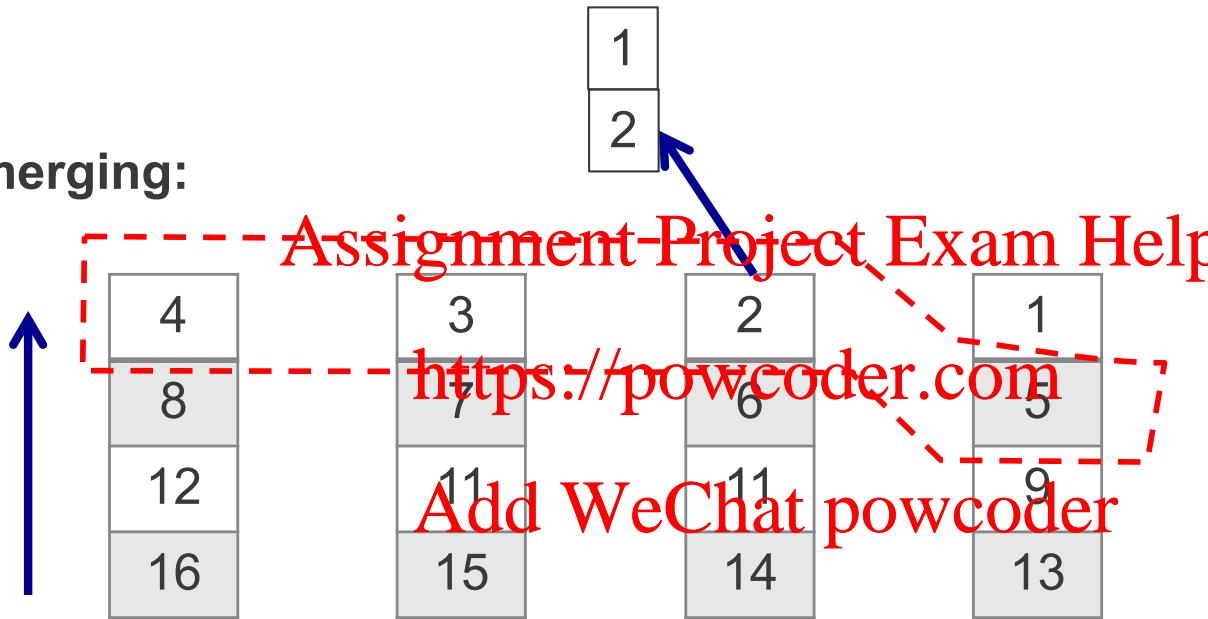
***k-way merging:***



Searching for the smallest value among the  $k$  values

## 4.3. Parallel External Sort (cont'd)

*k*-way merging:



Searching for the smallest value among the  $k$  values

and so on...

## 4.3. Parallel External Sort (cont'd)

- Parallel Binary-Merge Sort (Binary Merging step)

- Binary merging vs.  $k$ -way merging
- In **binary merging**, it is pairwise, but can be time consuming if the list is long
- System requirements: the pipeline process in binary merging requires extra overheads

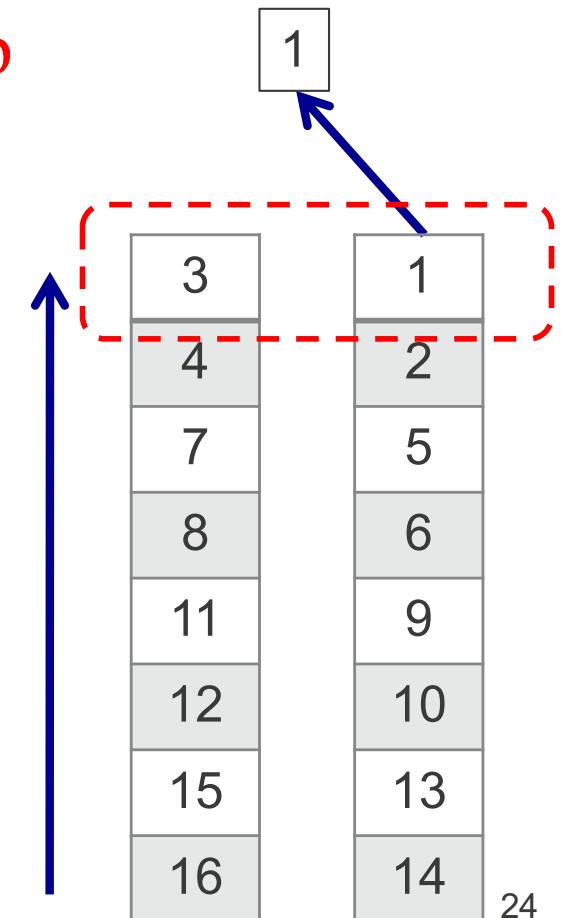
Assignment Project Exam Help

<https://powcoder.com>

Binary merging:

Add WeChat powcoder

Compare two values  
only, but lists are  
longer



## 4.3. Parallel External Sort (cont'd)

- Parallel Binary-Merge Sort (Binary Merging step)

- Binary merging vs.  $k$ -way merging
- In **binary merging**, it is pairwise, but can be time consuming if the list is long
- System requirements: the pipeline process in binary merging requires extra overheads

Assignment Project Exam Help

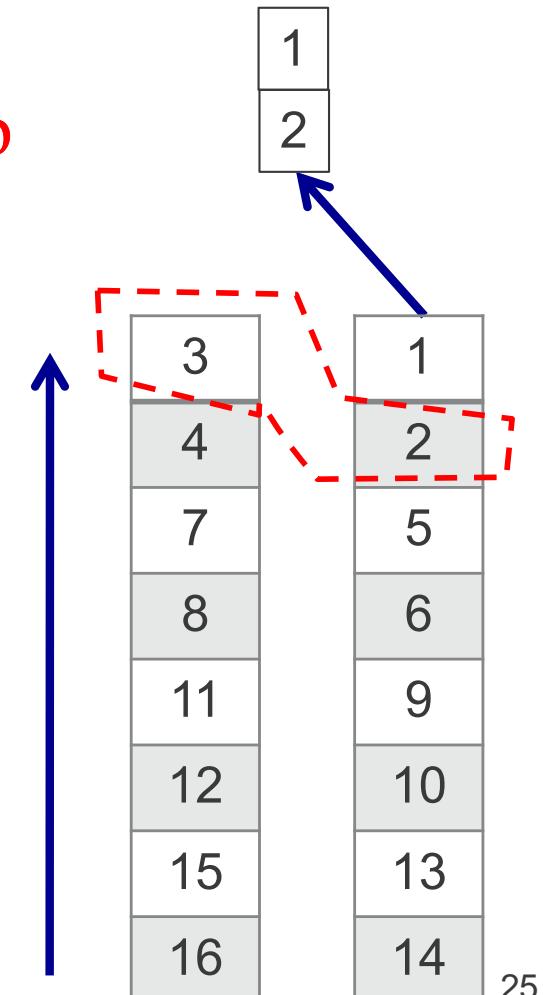
<https://powcoder.com>

Binary merging:

Add WeChat powcoder

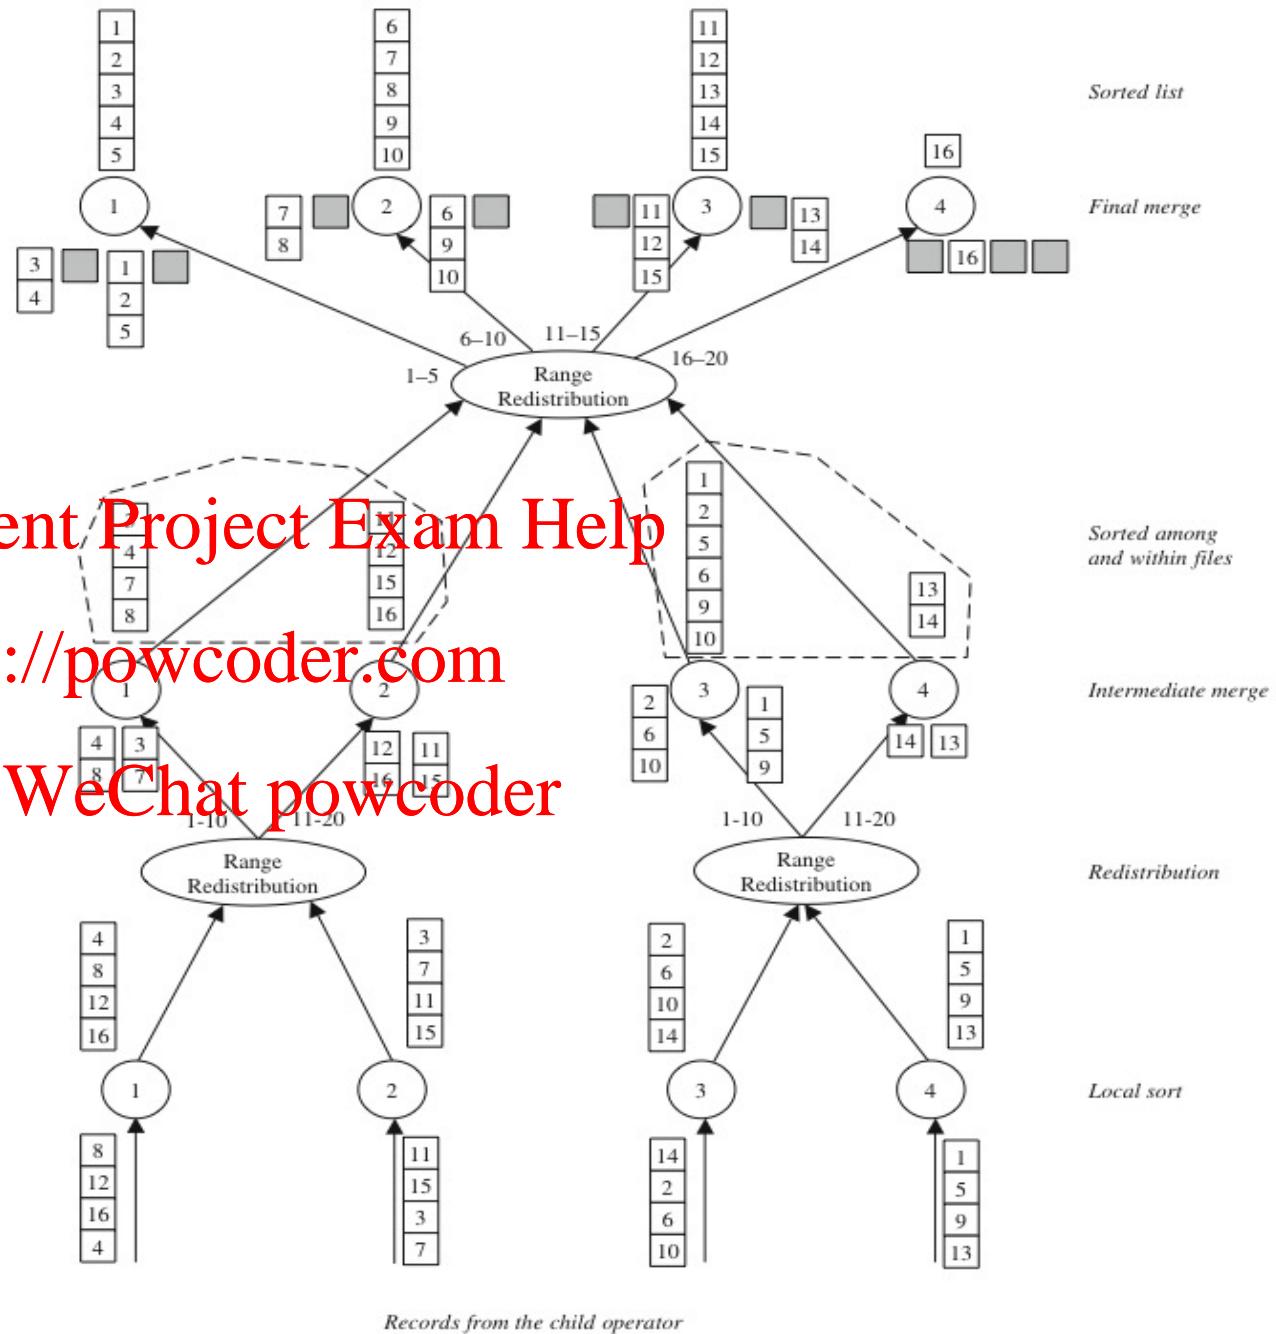
Compare two values only, but lists are longer

And so on...



## Parallel Redistribution Binary-Merge Sort

- Parallelism at all levels in the pipeline hierarchy
- Step 1: local sort
- Step 2: redistribute the results of local sort
- Step 3: merge using the same pool of processors
- Benefit: merging becomes lighter than without redistribution
- Problem: height of the tree



**Figure 4.6** Parallel redistribution binary-merge sort

## Parallel Redistribution Merge-All Sort

- Reduce the height of the tree, and still maintain parallelism
- Like parallel merge-all sort, but with redistribution
- The advantage is true parallelism in merging
- Skew problem in the merging

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

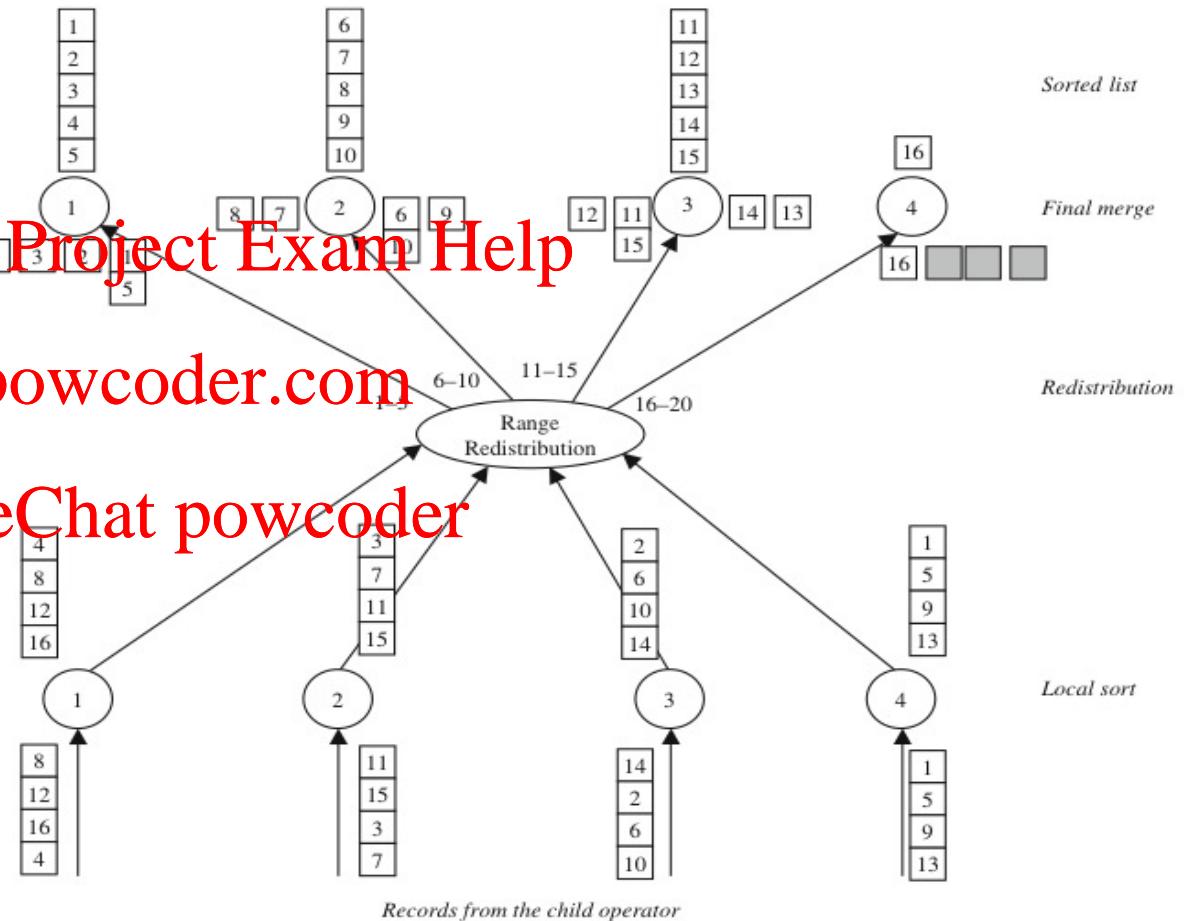


Figure 4.7 Parallel redistribution merge-all sort

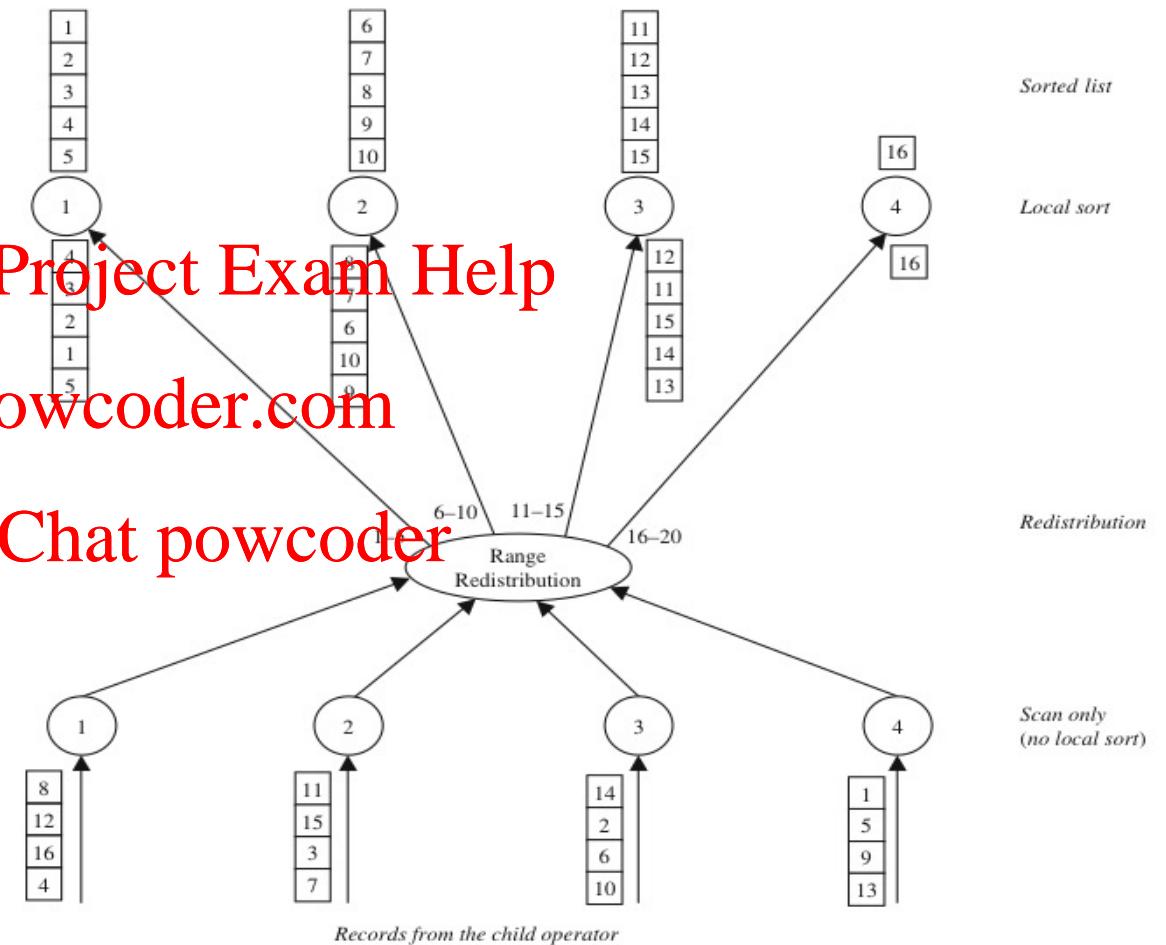
## • Parallel Partitioned Sort

- Two stages: Partitioning stage and Independent local work
- Partitioning (or range redistribution) may raise load skew
- Local sort is done after the partitioning, not before
- No merging is necessary
- Main problem: **Skew** produced by the partitioning

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



**Figure 4.8** Parallel partitioned sort

## 4.2. Serial External Sorting (cont'd)

- **Exercise 4 (Home Work)**

- Given a data set  $D = \{55; 30; 68; 39; 1; 4; 49; 90; 34; 76; 82; 56; 31; 25; 78; 56; 38; 32; 88; 9; 44; 98; 11; 70; 66; 89; 99; 22; 23; 26\}$  and four processors, show step by step how the Parallel Partitioned Sort works

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## 4.2. Serial External Sorting (cont'd)

- **Exercise 5 (Difficult)**

- Given the same dataset as in the previous question, and 4 processors, show how **Load Balancing** is achieved in the **Parallel Partitioned Sort**.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# FIT5202 (Volume IV – Sort and Group By) Assignment Project Exam Help

Week 4b – Parallel Group By

<https://powcoder.com>

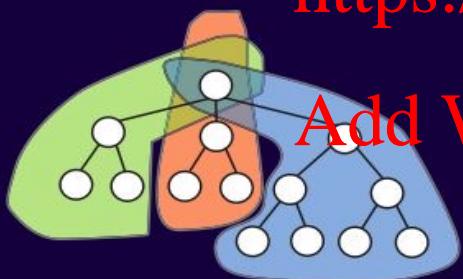
**algorithm** distributed systems **database**  
systems **computation** knowledge management  
**design** e-business **model** data mining **int**  
distributed systems **database** software  
**computation** knowledge management **an**

TANIAR  
LEUNG  
RAHAYU  
GOEL

*Wiley Series on Parallel and Distributed Computing • Albert Zomaya, Series Editor*

High Performance Parallel  
Database Processing and  
Grid Databases

High Performance Parallel Database  
Processing and Grid Databases



DAVID TANIAR, CLEMENT H.C. LEUNG,  
WENNY RAHAYU, and SUSHANT GOEL



WILEY

# Chapter 4

## Parallel Sort and GroupBy

<https://powcoder.com>

Add WeChat powcoder

- 4.1 Sorting, Duplicate Removal and Aggregate
- 4.2 Serial External Sorting Method
- 4.3 Algorithms for Parallel External Sort
- 4.4 Parallel Algorithms for GroupBy Queries
- 4.5 Cost Models for Parallel Sort
- 4.6 Cost Models for Parallel GroupBy
- 4.7 Summary
- 4.8 Bibliographical Notes
- 4.9 Exercises

## 4.1. GroupBy, and Serial GroupBy

Select Suburb, Count(\*)  
From Student  
Group By Suburb;

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Erik	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## 4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(\*)  
From Student  
Group By Suburb;

### Processing Steps:

1. Read the first student record, and hash the suburb to the hash table

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Erik	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash Table

1		
2		
3		
4		
5		
6		
7		
8	Clayton	1
9		

Hash the record using a certain hash function

## 4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(\*)  
From Student  
Group By Suburb;

### Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Erik	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn



Hash Table

1	Hawthorn	1
2		
3		
4		
5		
6		
7		
8	Clayton	1
9		

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## 4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(\*)  
From Student  
Group By Suburb;

### Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it
3. Read the subsequent records one-by-one and hash them

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Erik	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash Table

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## 4.1. Serial GroupBy Processing (cont'd)

```
Select Suburb, Count(*)  
From Student  
Group By Suburb;
```

### Processing Steps:

1. Read the first student record, and hash the suburb to the hash table
2. Read the second record and hash it
3. Read the subsequent records one-by-one and hash them
4. Read the Hash Table, and store this in disk as the query results

Query Results in Disk

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

Hawthorn	3
Caulfield	2
Malvern	1
Balwyn	1
Kew	1
Richmond	1
Elwood	1
Clayton	2
Doncaster	1

Hash Table in Main-Memory

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

## 4.1. Serial GroupBy Processing (cont'd)

Select Suburb, Count(\*)  
From Student  
Group By Suburb;

Student	Suburb
Adam	Clayton
Ben	Hawthorn

This will work if we assume that the main-memory can hold the entire Hash Table.

**Assignment Project Exam Help**  
<https://powcoder.com>

How about if the Hash Table is so big that it cannot fit into the main memory?

For example, how about if the main-memory can only hold 4 hash records at a time? How does the Group By processing work?

Leonard	Balwyn
Mary	Hawthorn

**Hash Table**

1	Hawthorn	3
2	Caulfield	2
3	Malvern	1
4	Balwyn	1
5	Kew	1
6	Richmond	1
7	Elwood	1
8	Clayton	2
9	Doncaster	1

## 4.1. Serial GroupBy Processing (cont'd)

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash Table


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assume that the main-memory can hold 4 records in the hash table.

It needs a bigger hash table, but it doesn't have.

Student	Suburb
Adam	Clayton
Ben	Hawthorn
Chris	Doncaster
Daniel	Caulfield
Eric	Kew
Fred	Richmond
Garry	Hawthorn
Harold	Elwood
Irene	Clayton
Jessica	Caulfield
Katie	Malvern
Leonard	Balwyn
Mary	Hawthorn

Hash Data  
Partitioning  
based on the  
**Suburb**

**Assignment Project Exam Help**

<https://powcoder.com>

Add WeChat **powcoder**

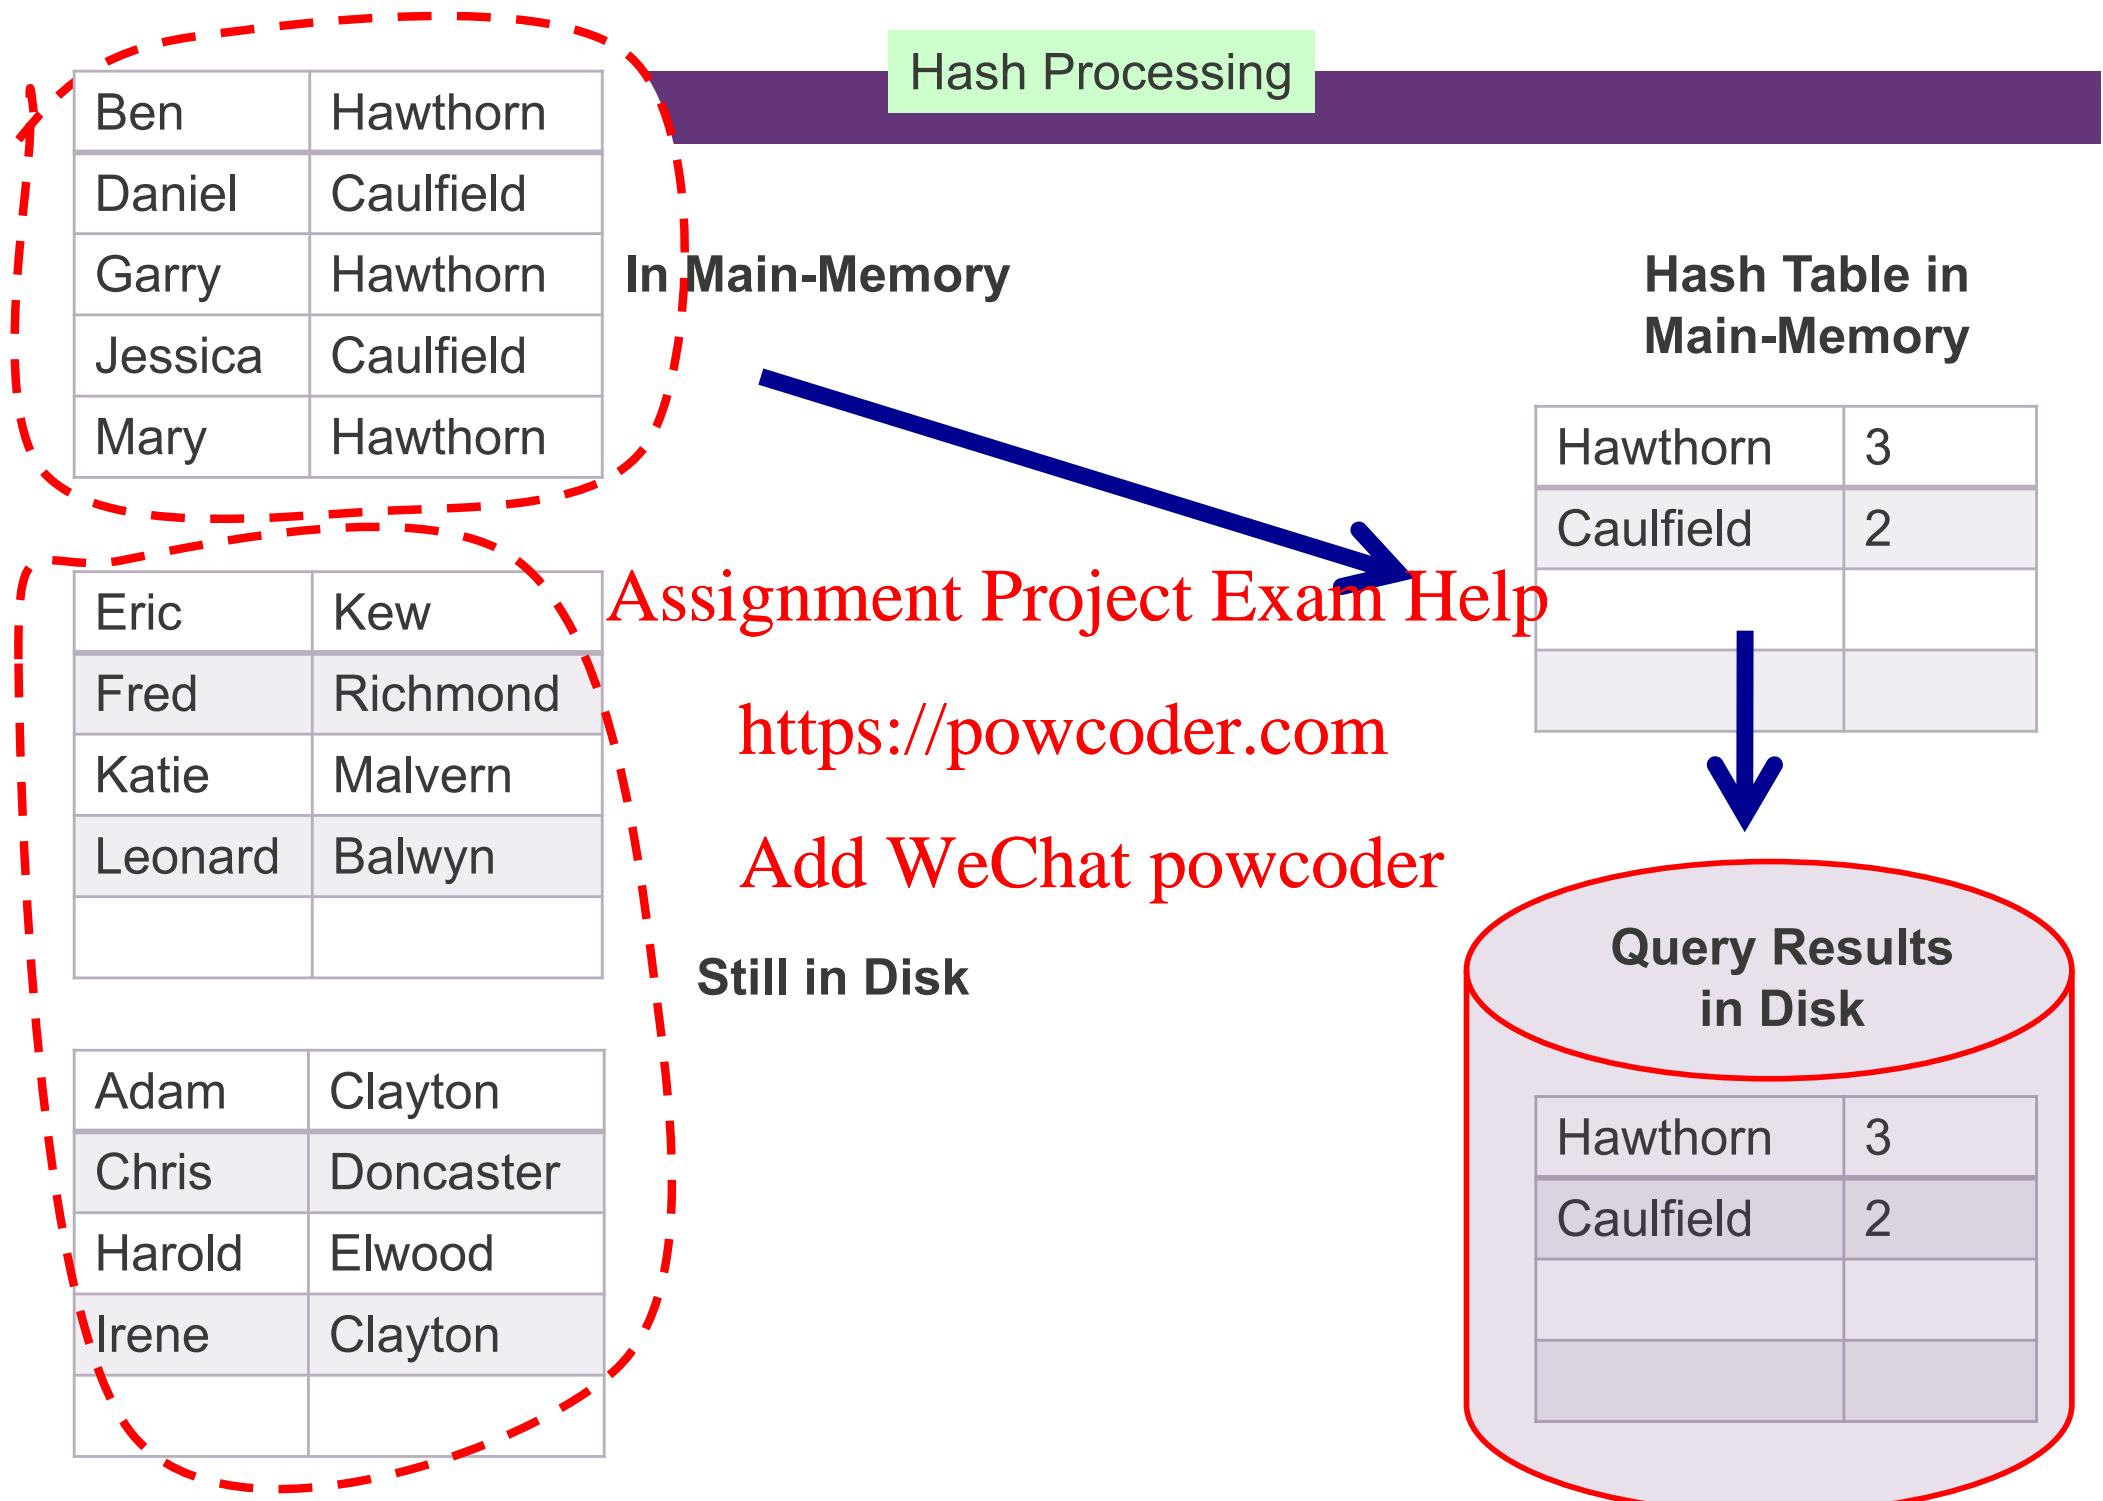
Ben	Hawthorn
Daniel	Caulfield
Garry	Hawthorn
Jessica	Caulfield
Mary	Hawthorn

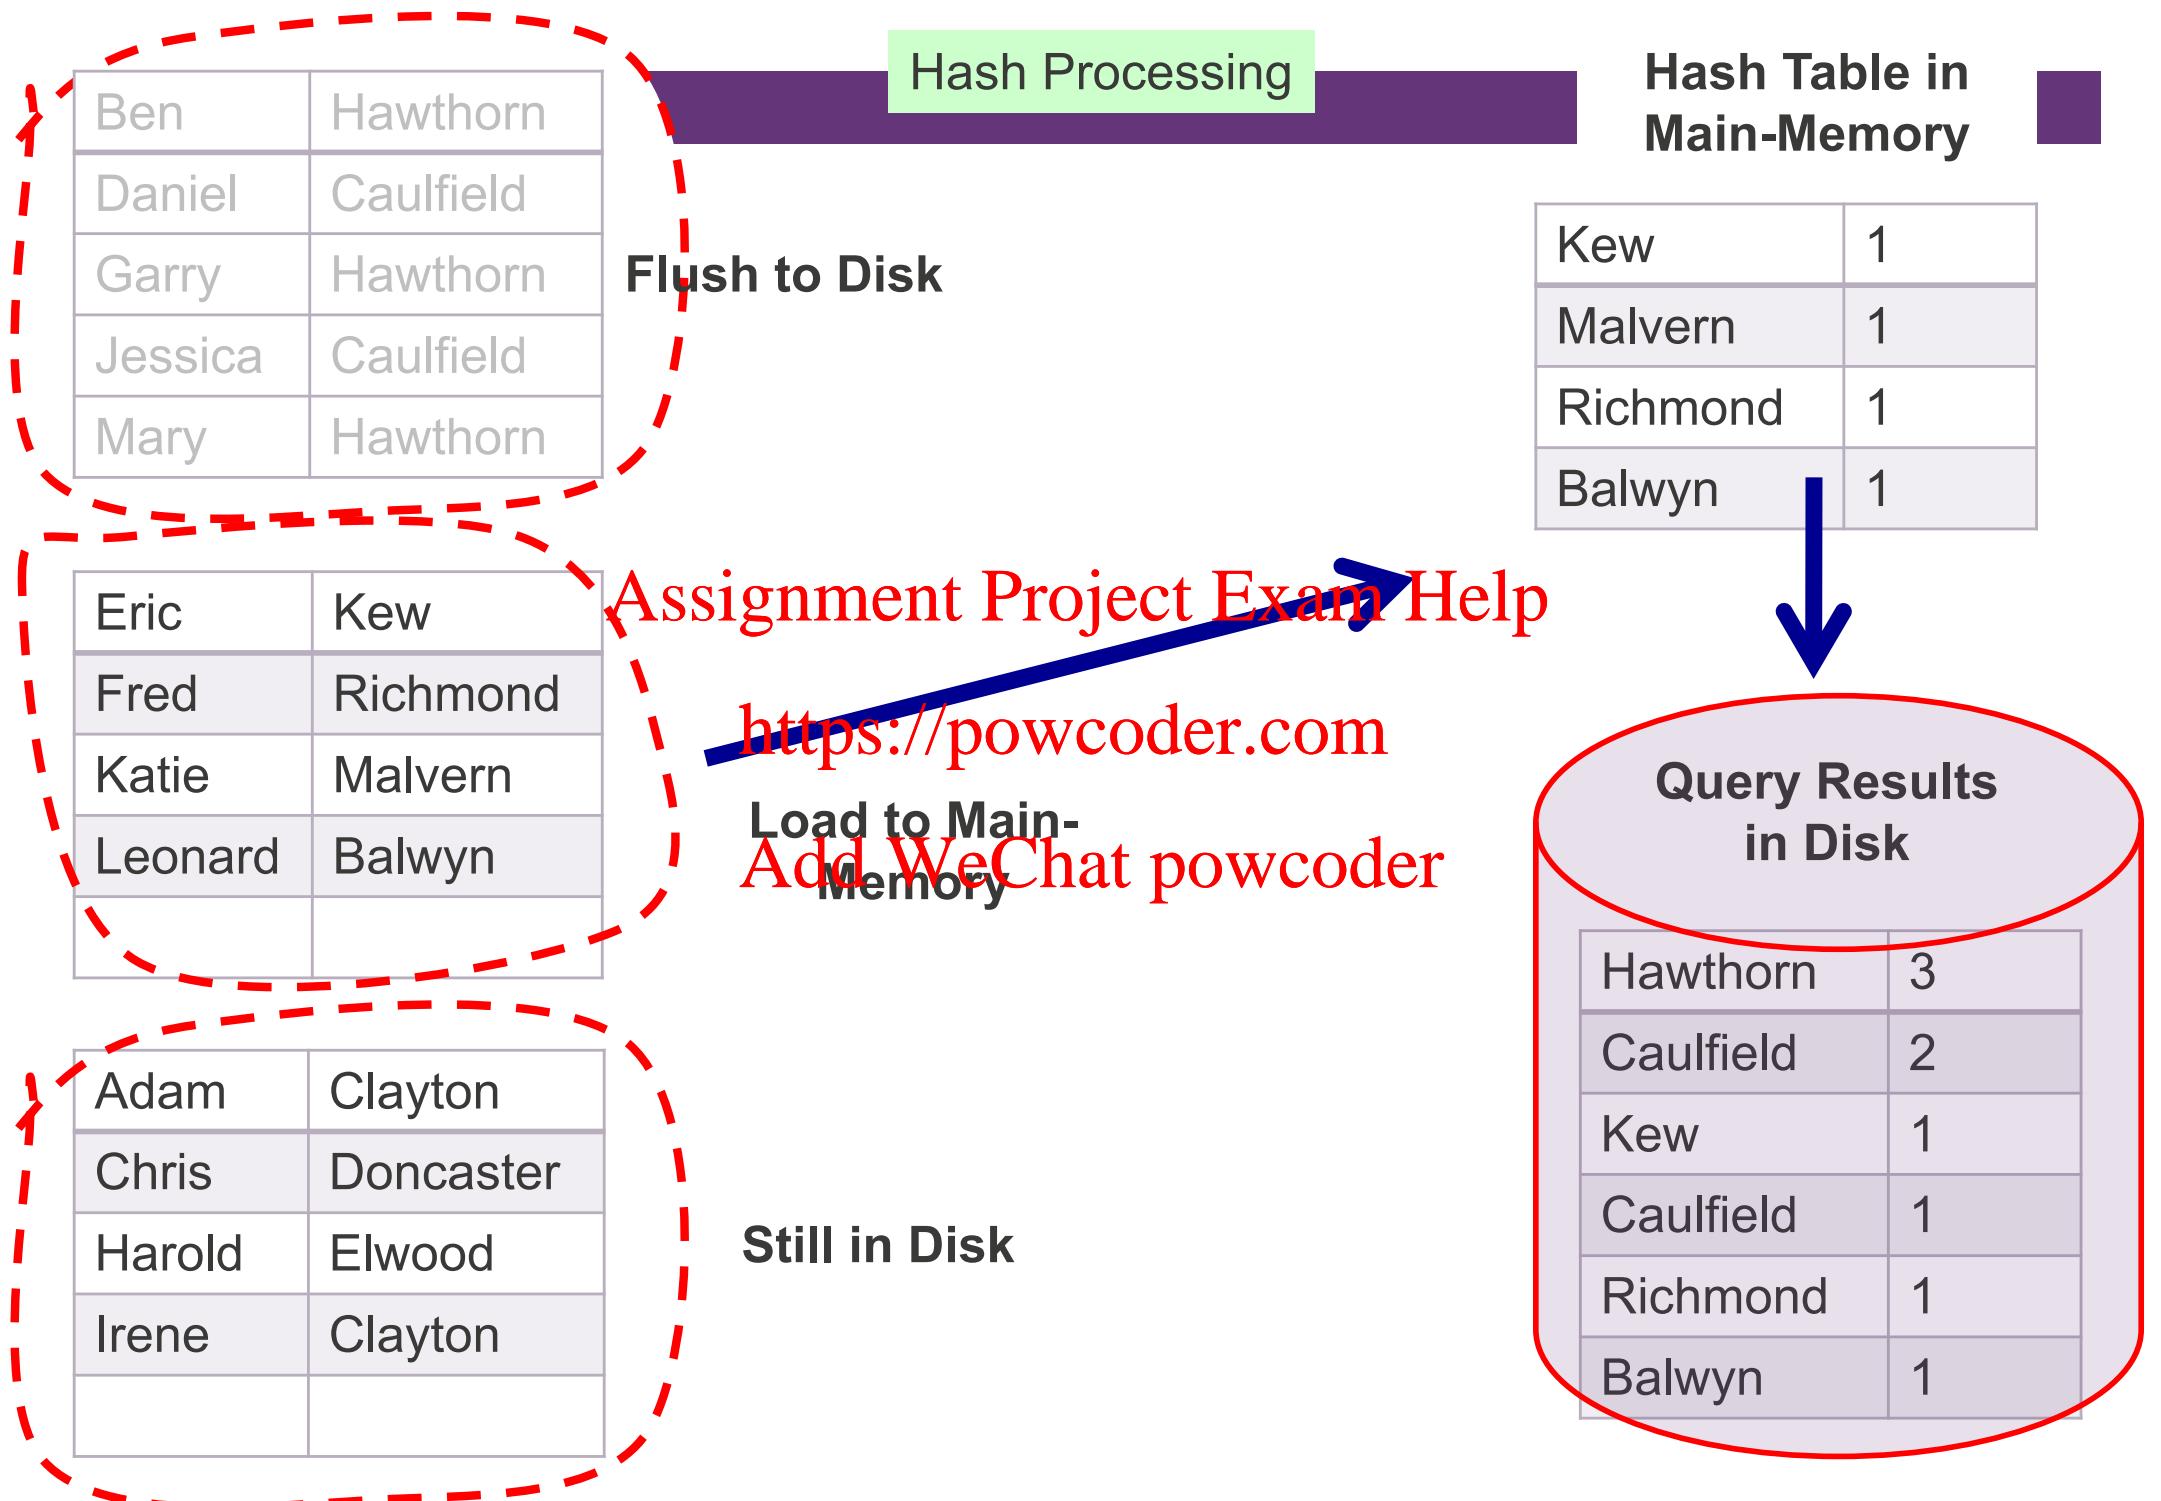
Eric	Kew
Fred	Richmond
Katie	Malvern
Leonard	Balwyn

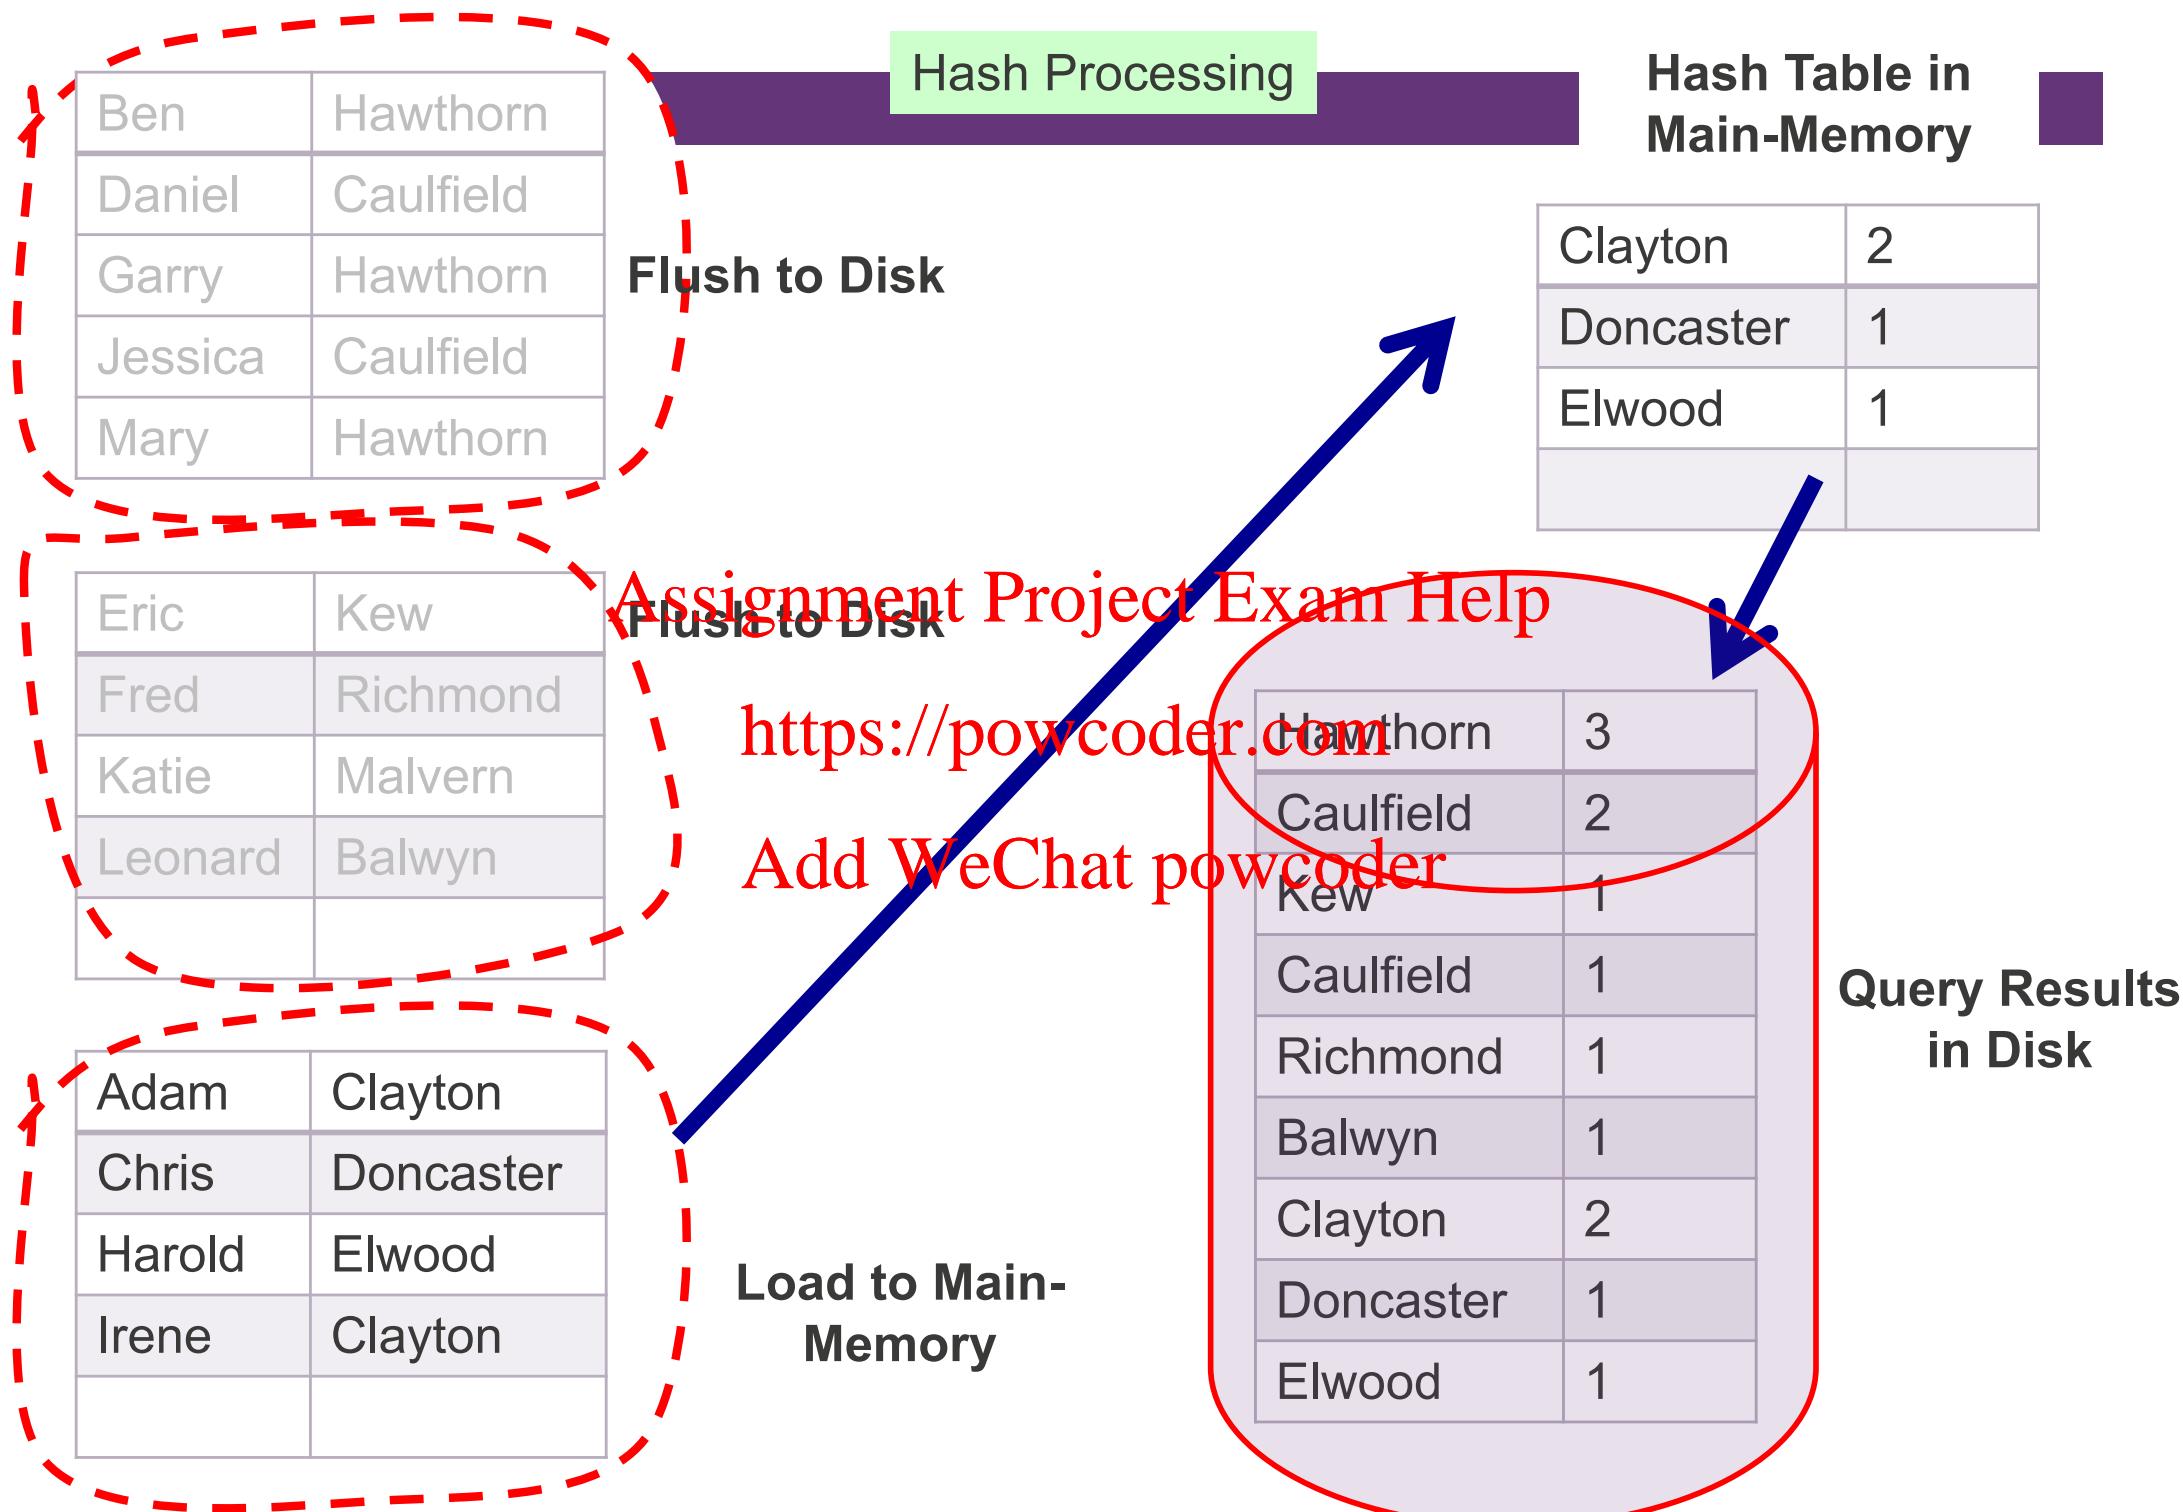
Adam	Clayton
Chris	Doncaster
Harold	Elwood
Irene	Clayton

In Main-Memory

In Disk







## 4.4. Parallel GroupBy

- Traditional methods (Merge-All and Hierarchical Merging)
- Two-phase method
- Redistribution methods

Without data  
redistribution

With data  
redistribution

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## 4.4. Parallel GroupBy (cont'd)

### • Traditional Methods

- Step 1: local aggregate in each processor
- Step 2: global aggregation
- May use a Merge-All or Hierarchical method
- Need to pay a special attention to some aggregate functions (AVG) when performing a local aggregate process

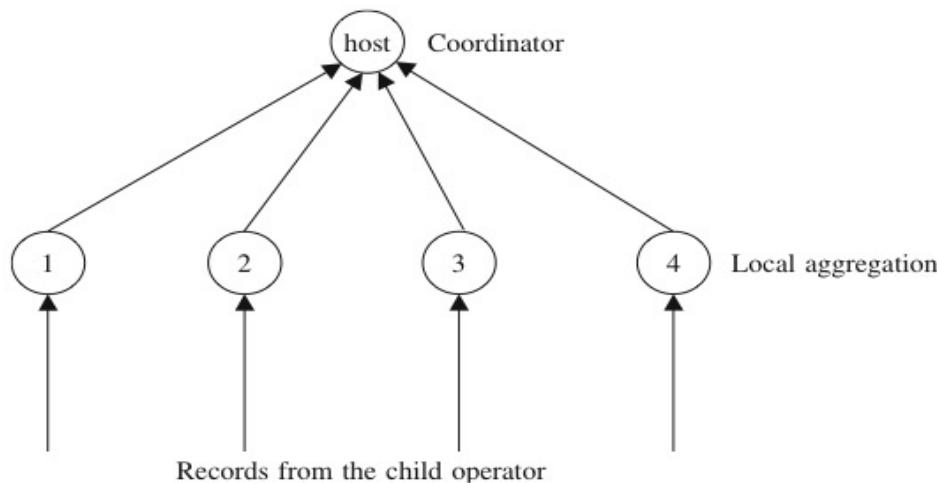


Figure 4.10 Traditional method

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

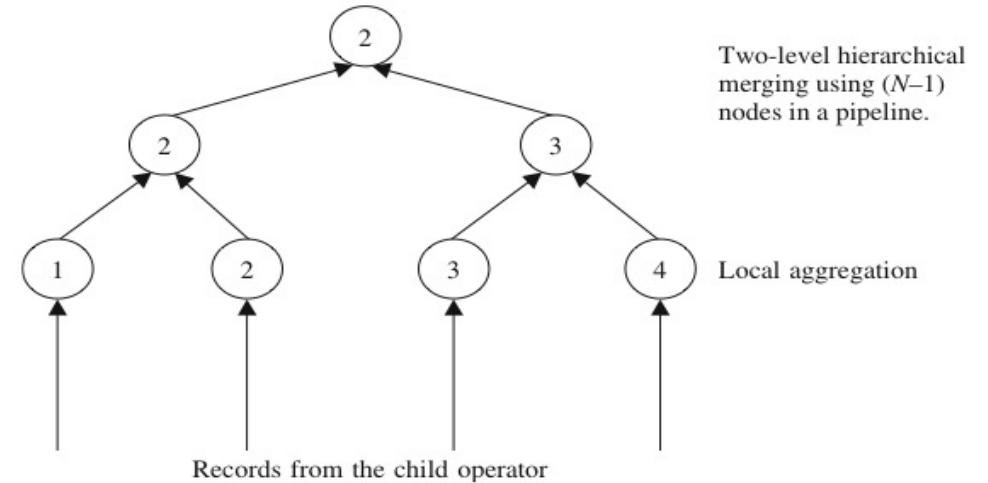
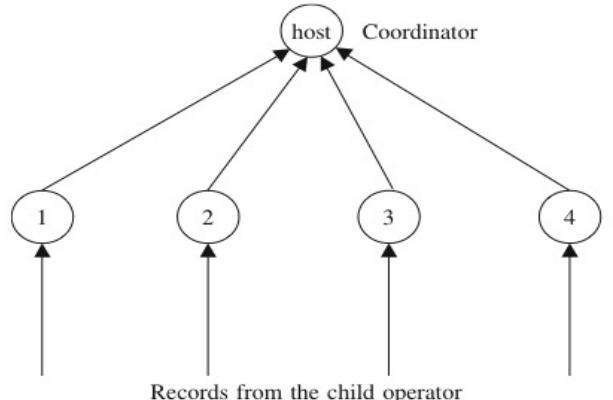


Figure 4.11 Hierarchical merging method

## 4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1



Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Local Aggregation Phase  
<https://powcoder.com>  
Add WeChat powcoder

Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

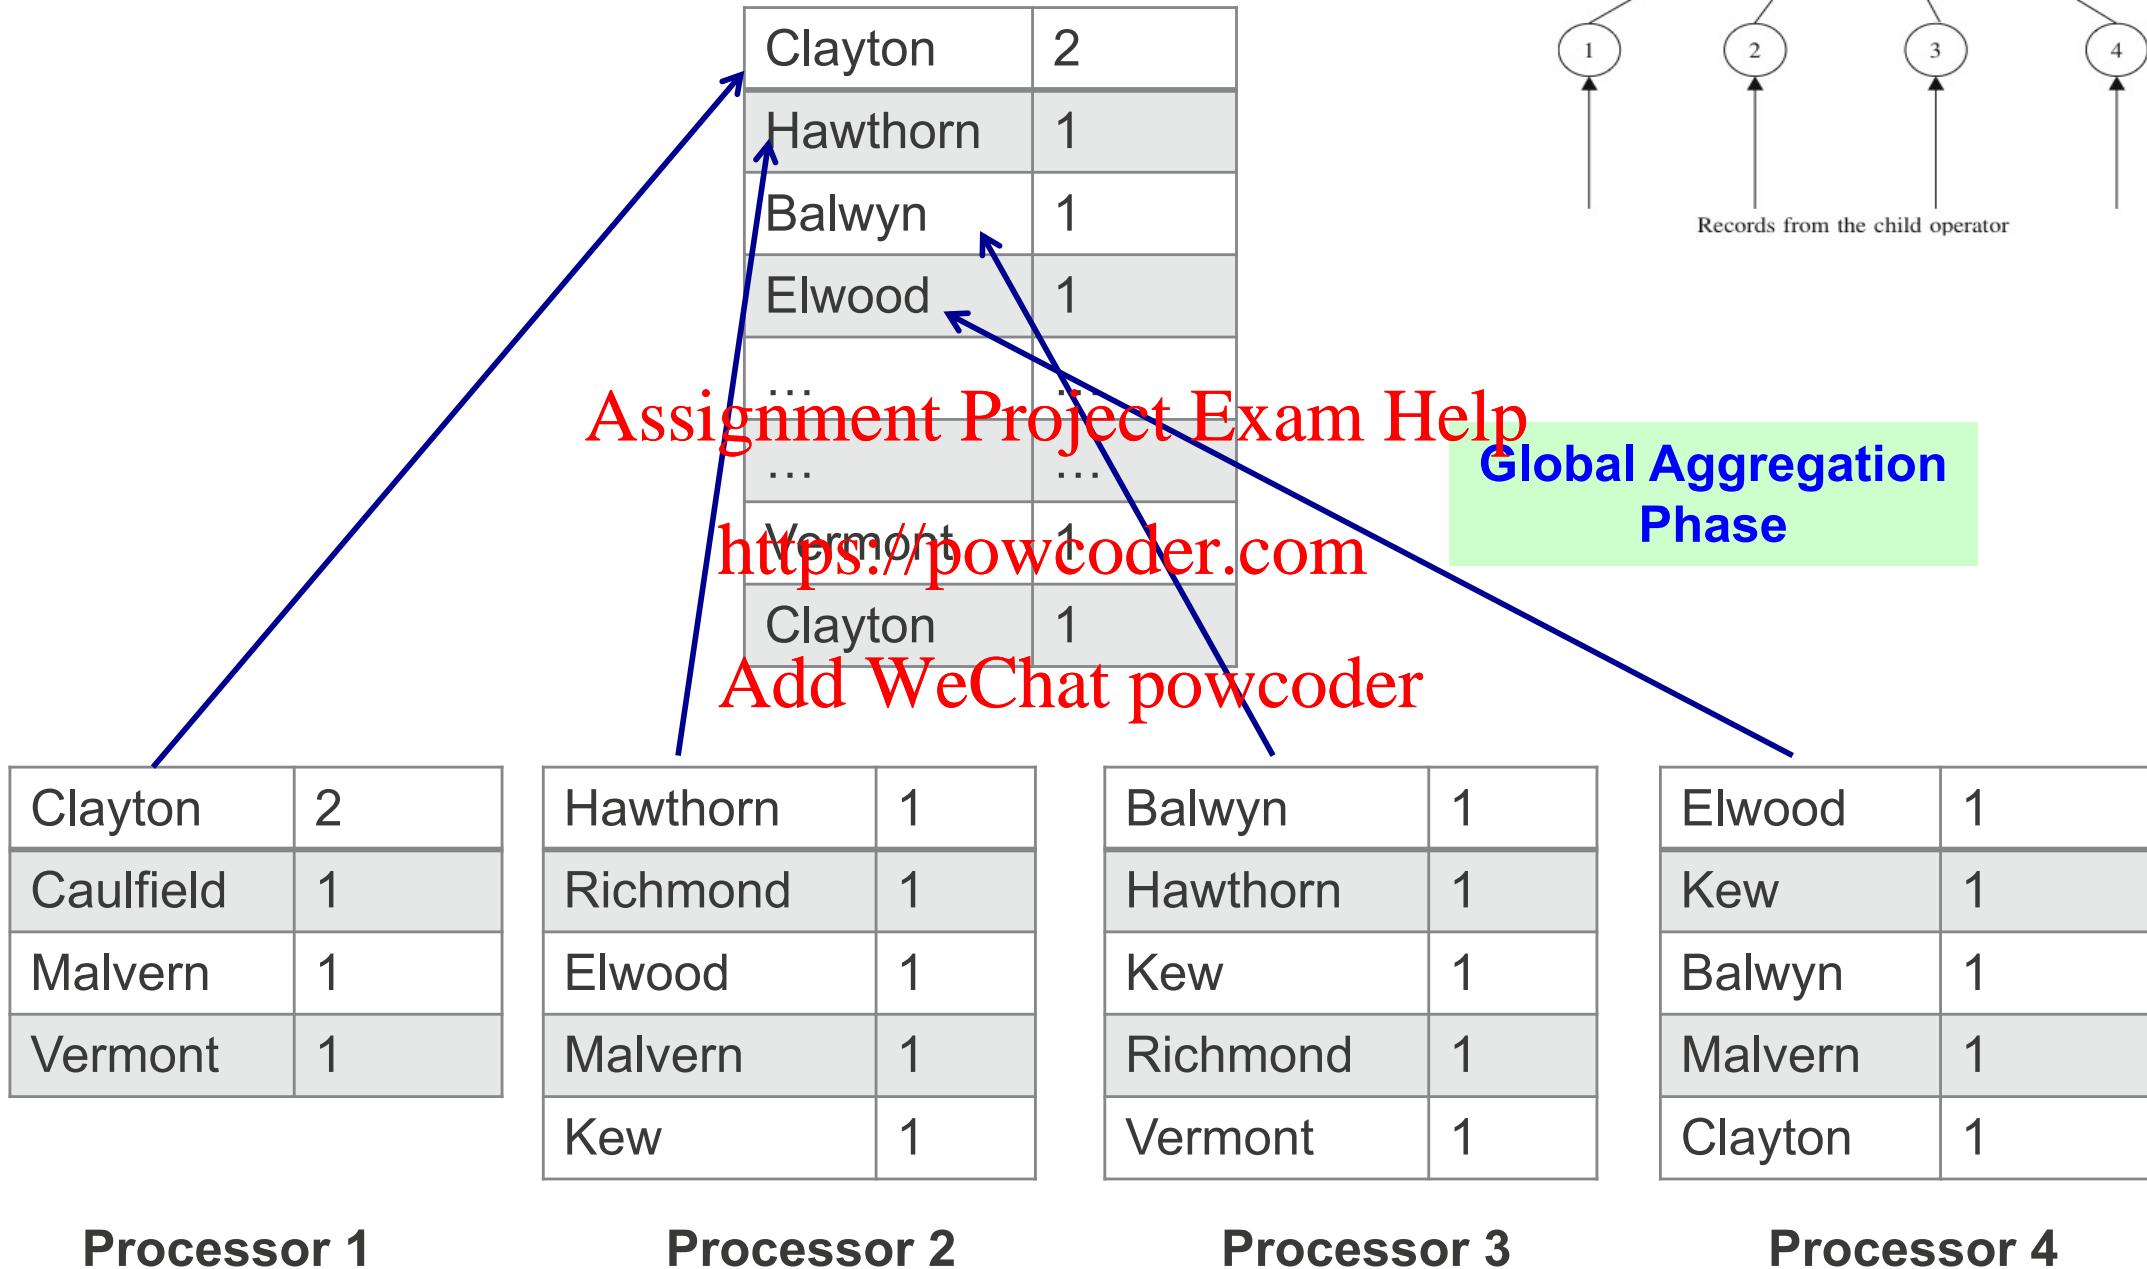


Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

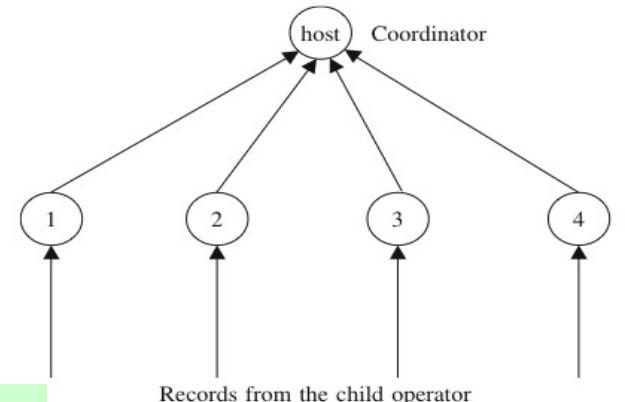


## 4.4. Parallel GroupBy (cont'd)

- Traditional Method: Merge All

Clayton	3
Hawthorn	2
Balwyn	2
...	...
Vermont	2

Final results



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Clayton	2
Hawthorn	1
Balwyn	1
Elwood	1
...	...
...	...
Vermont	1
Clayton	1

Global Aggregation Phase

## **• Exercise 6 (FLUX Quiz)**

- The limitations of the Traditional Approach (Merge All) to process a Group By query are:
  - A. Global aggregation is carried out by one processor
  - B. Network bottleneck when sending the local aggregation results to the coordinator
  - C. No parallelism in the global aggregation phase
  - D. All of the above
  - E. Some of the above

**Assignment Project Exam Help**

**<https://powcoder.com>**

**Add WeChat powcoder**

## 4.4. Parallel GroupBy (cont'd)

- **Two-Phase Method**

- Step 1: local aggregate in each processor. Each processor groups local records according to the groupby attribute
- Step 2: global aggregation where all temp results from each processor are redistributed and then final aggregate is performed in each processor

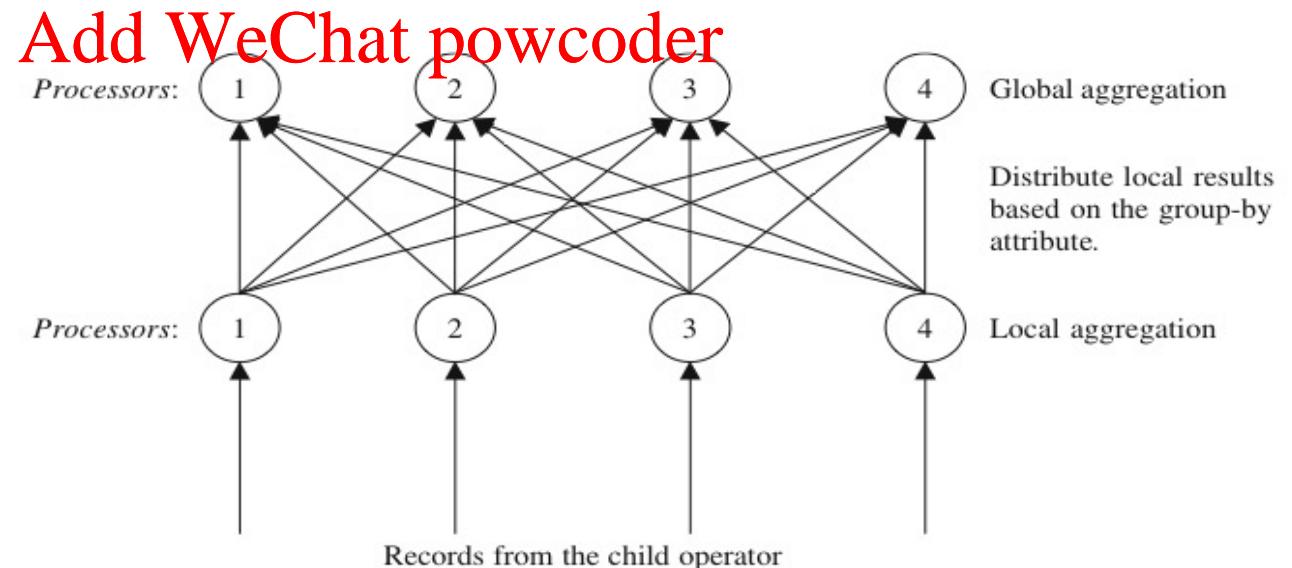
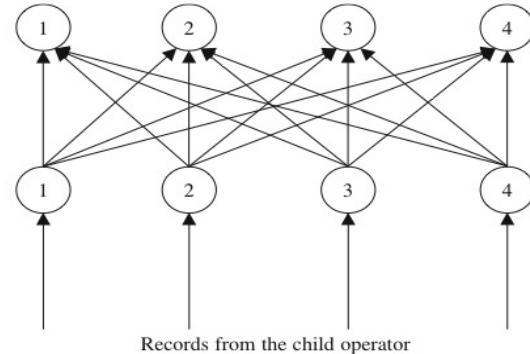


Figure 4.12 Two-phase method

## 4.4. Parallel GroupBy (cont'd)

- Two-Phase Method



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1



Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Local Aggregation Phase  
<https://powcoder.com>  
Add WeChat powcoder

Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont



Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

Clayton	2
Balwyn	1
Elwood	1
Caulfield	1
Elwood	1
Balwyn	1
Clayton	1

Hawthorn	1
Hawthorn	1
Kew	1
Kew	1
Kew	1

Malvern	1
Malvern	1
Malvern	1

Richmond	1
Vermont	1
Richmond	1
Vermont	1

Assignment Project Exam Help

Distribute Local Aggregation  
Results Phase

<https://powcoder.com>

Add WeChat powcoder

Clayton	2
Caulfield	1
Malvern	1
Vermont	1

Hawthorn	1
Richmond	1
Elwood	1
Malvern	1
Kew	1

Balwyn	1
Hawthorn	1
Kew	1
Richmond	1
Vermont	1

Elwood	1
Kew	1
Balwyn	1
Malvern	1
Clayton	1

Processor 1

Processor 2

Processor 3

Processor 4

## 4.4. Parallel GroupBy (cont'd)

- Two-Phase Method

Final results

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1

Hawthorn	2
Kew	3

Malvern	3
---------	---

Richmond	2
Vermont	2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder  
Global Aggregation Phase

Clayton	2
Balwyn	1
Elwood	1
Caulfield	1
Elwood	1
Balwyn	1
Clayton	1

Hawthorn	1
Hawthorn	1
Kew	1
Kew	1
Kew	1

Malvern	1
Malvern	1
Malvern	1

Richmond	1
Vermont	1
Richmond	1
Vermont	1

Processor 1

Processor 2

Processor 3

Processor 4

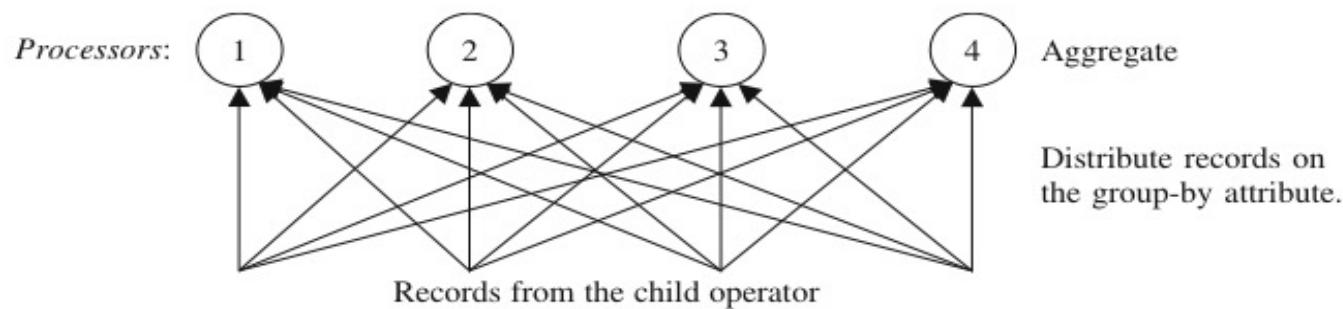
## 4.4. Parallel GroupBy (cont'd)

- **Redistribution Method**

- Step 1 (Partitioning phase): redistribute raw records to all processors
- Step 2 (Aggregation phase): each processor performs a local aggregation

<https://powcoder.com>

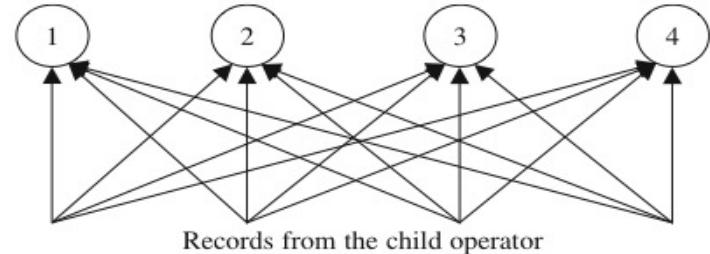
Add WeChat powcoder



**Figure 4.13** Redistribution method

## 4.4. Parallel GroupBy (cont'd)

- Redistribution Method



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Processor 2

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Redistribution Method

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

Processor 1

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Partitioning Phase

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Processor 2

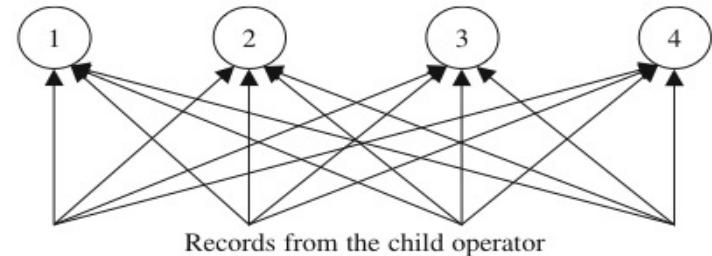
Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Processor 3

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Processor 4

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton



## 4.4. Parallel GroupBy (cont'd)

### • Redistribution Method

Final results

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1

Hawthorn	2
Kew	3

Malvern	3
---------	---

Richmond	2
Vermont	2

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 1

Processor 2

Processor 3

Processor 4

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder  
Aggregation Phase

## 4.4. Parallel GroupBy (cont'd)

### • Redistribution Method

Clayton	3
Balwyn	2
Elwood	2
Caulfield	1

Hawthorn	2
Kew	3

Malvern	3

Richmond	2
Vermont	2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder  
What is the problem here?

Adam	Clayton
Kelly	Balwyn
Peter	Elwood
Ben	Clayton
Chris	Caulfield
Harold	Elwood
Roger	Balwyn
Tracy	Clayton

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 1

Processor 2

Processor 3

Processor 4

## 4.4. Parallel GroupBy (cont'd)

- Redistribution Method (Task Stealing)

Adam	Clayton
Kelly	Balwyn
Ben	Clayton
Chris	Caulfield
Roger	Balwyn
Tracy	Clayton
Peter	Elwood
Harold	Elwood

Processor 1

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Processor 2

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

Processor 3

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

Adam	Clayton
Ben	Clayton
Chris	Caulfield
Dennis	Malvern
Eric	Vermont

Fred	Hawthorn
George	Richmond
Harold	Elwood
Irene	Malvern
Jessica	Kew

Kelly	Balwyn
Lesley	Hawthorn
Megan	Kew
Naomi	Richmond
Oscar	Vermont

Peter	Elwood
Quin	Kew
Roger	Balwyn
Sarah	Malvern
Tracy	Clayton

## 4.4. Parallel GroupBy (cont'd)

- Redistribution Method (Task Stealing)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Adam	Clayton
Kelly	Balwyn
Ben	Clayton
Chris	Caulfield
Roger	Balwyn
Tracy	Clayton

Peter	Elwood
Harold	Elwood

Processor 1

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Processor 2

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

Task stealing

Peter	Elwood
Harold	Elwood

Processor 3

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

## 4.4. Parallel GroupBy (cont'd)

### • Redistribution Method (Task Stealing)

Clayton	3
Balwyn	2
Caulfield	1



Hawthorn	2
Kew	3

Malvern	3
Elwood	2

Richmond	2
Vermont	2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Adam	Clayton
Kelly	Balwyn
Ben	Clayton
Chris	Caulfield
Roger	Balwyn
Tracy	Clayton

Processor 1

Fred	Hawthorn
Lesley	Hawthorn
Quin	Kew
Megan	Kew
Jessica	Kew

Processor 2

Dennis	Malvern
Irene	Malvern
Sarah	Malvern

Peter	Elwood
Harold	Elwood

Processor 3

George	Richmond
Naomi	Richmond
Eric	Vermont
Oscar	Vermont

Processor 4

## **• Exercise 7 (FLUX Quiz)**

- The Redistribution Method has a load balancing option, through the Task Stealing method. The Two-Phase Method does not have a load balancing problem.

**Assignment Project Exam Help**

- A. TRUE
- B. FALSE

**<https://powcoder.com>**

**Add WeChat powcoder**

## 4.7. Summary

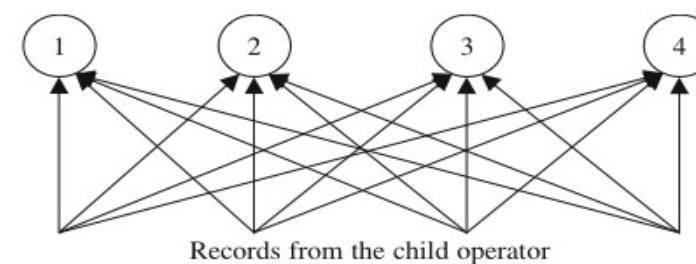
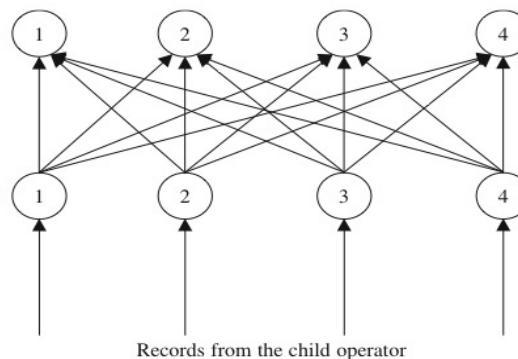
- Parallel groupby algorithms
  - Traditional methods (merge-all and hierarchical methods)
  - Two-phase method - – Local aggregation **before** data redistribution
  - Redistribution method - Local aggregation **after** data redistribution
- **Two-phase** and **Redistribution** methods perform better than the traditional and hierarchical merging methods

Add WeChat powcoder

- **Two-phase method** works well when the number of groups is small, whereas the **Redistribution method** works well when the number of groups is large

Ambuj and Naughton. "Adaptive parallel aggregation algorithms." (1995):

Why??



## • Homework Exercises

- 1. Show how Load Balancing through Task Stealing be achieved in the Two Phase Method (using the same sample data as above) – EASY

**Assignment Project Exam Help**

- 2. Why is the Two-Phase Method good when the number of groups is small, whereas the Redistribution Method good when the number of groups is large? – MORE CHALLENGING  
<https://powcoder.com>
- 3. In what scenario may super linear speedup be achieved? – MORE CHALLENGING  
(Hints: See slides #8-#13 → the hash table cannot fit into main-memory)